

Concurrent Systems Operating Systems

3D4 ← → CS2016

Andrew Butterfield
ORI.G39, Andrew.Butterfield@scss.tcd.ie



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

with thanks to Mike Brady

Overview

- **Concurrency**

- What / Why / How,
- Grappling with Concurrency Issues.

- **Operating Systems**

- Operating system architectures,
- Memory Management (OS perspective),
- Processes and Thread Management,
- File Storage (disk I/O and file systems).



Practical Matters — Linux

- Course is based on Linux and the standard C program build toolset, “Build Essentials”, POSIX Threads, SPIN, etc.
- Mac OS X also works pretty well.
- Windows not really supported or suitable.
- Suggest you consider running Linux in a Virtual Machine:
 - VMWare / VMWare Fusion for Mac — very good, widely used,
 - Virtual Box — free.
- Ubuntu / Debian. Don't go crazy.



Practical Matters — Labs

- 20% Labs/Practicals, 80% Examination
- Four Practicals, worth 2%, 6%, 6% and 6%.
 - First practical - simple exercise in compiling and running some concurrent code
 - Due: Monday Feb. 4th, 09:00 (via Blackboard)



Practical Matters — Software

- VMWare Academic and Microsoft Developer Licenses:
https://support.scss.tcd.ie/School_Site_Licences



POSIX

- POSIX – Portable Operating System Interface
 - for variants of Unix, including Linux
 - IEEE 1003, ISO/IEC9945
- Really, considered a standard set of facilities and APIs for Unix.
 - 1988 onwards
 - Doesn't *have* to be Unix – e.g. Cygwin for Windows give it partial POSIX compliance
 - ref: <http://en.wikipedia.org/wiki/POSIX>



POSIX Threads

- POSIX Threads aka '*pthread*'
 - correspond to 'Light Weight Processes' (LWPs) in older literature.
 - *pthread*s live within processes.
 - processes have separate memory spaces from one another
 - thus, inter-process communication & scheduling may be expensive
 - *pthread*s (within a process) share the same memory space
 - inter-thread communication & scheduling can be cheap
- *Classic tradeoff: speed vs. stability/ruggedness*

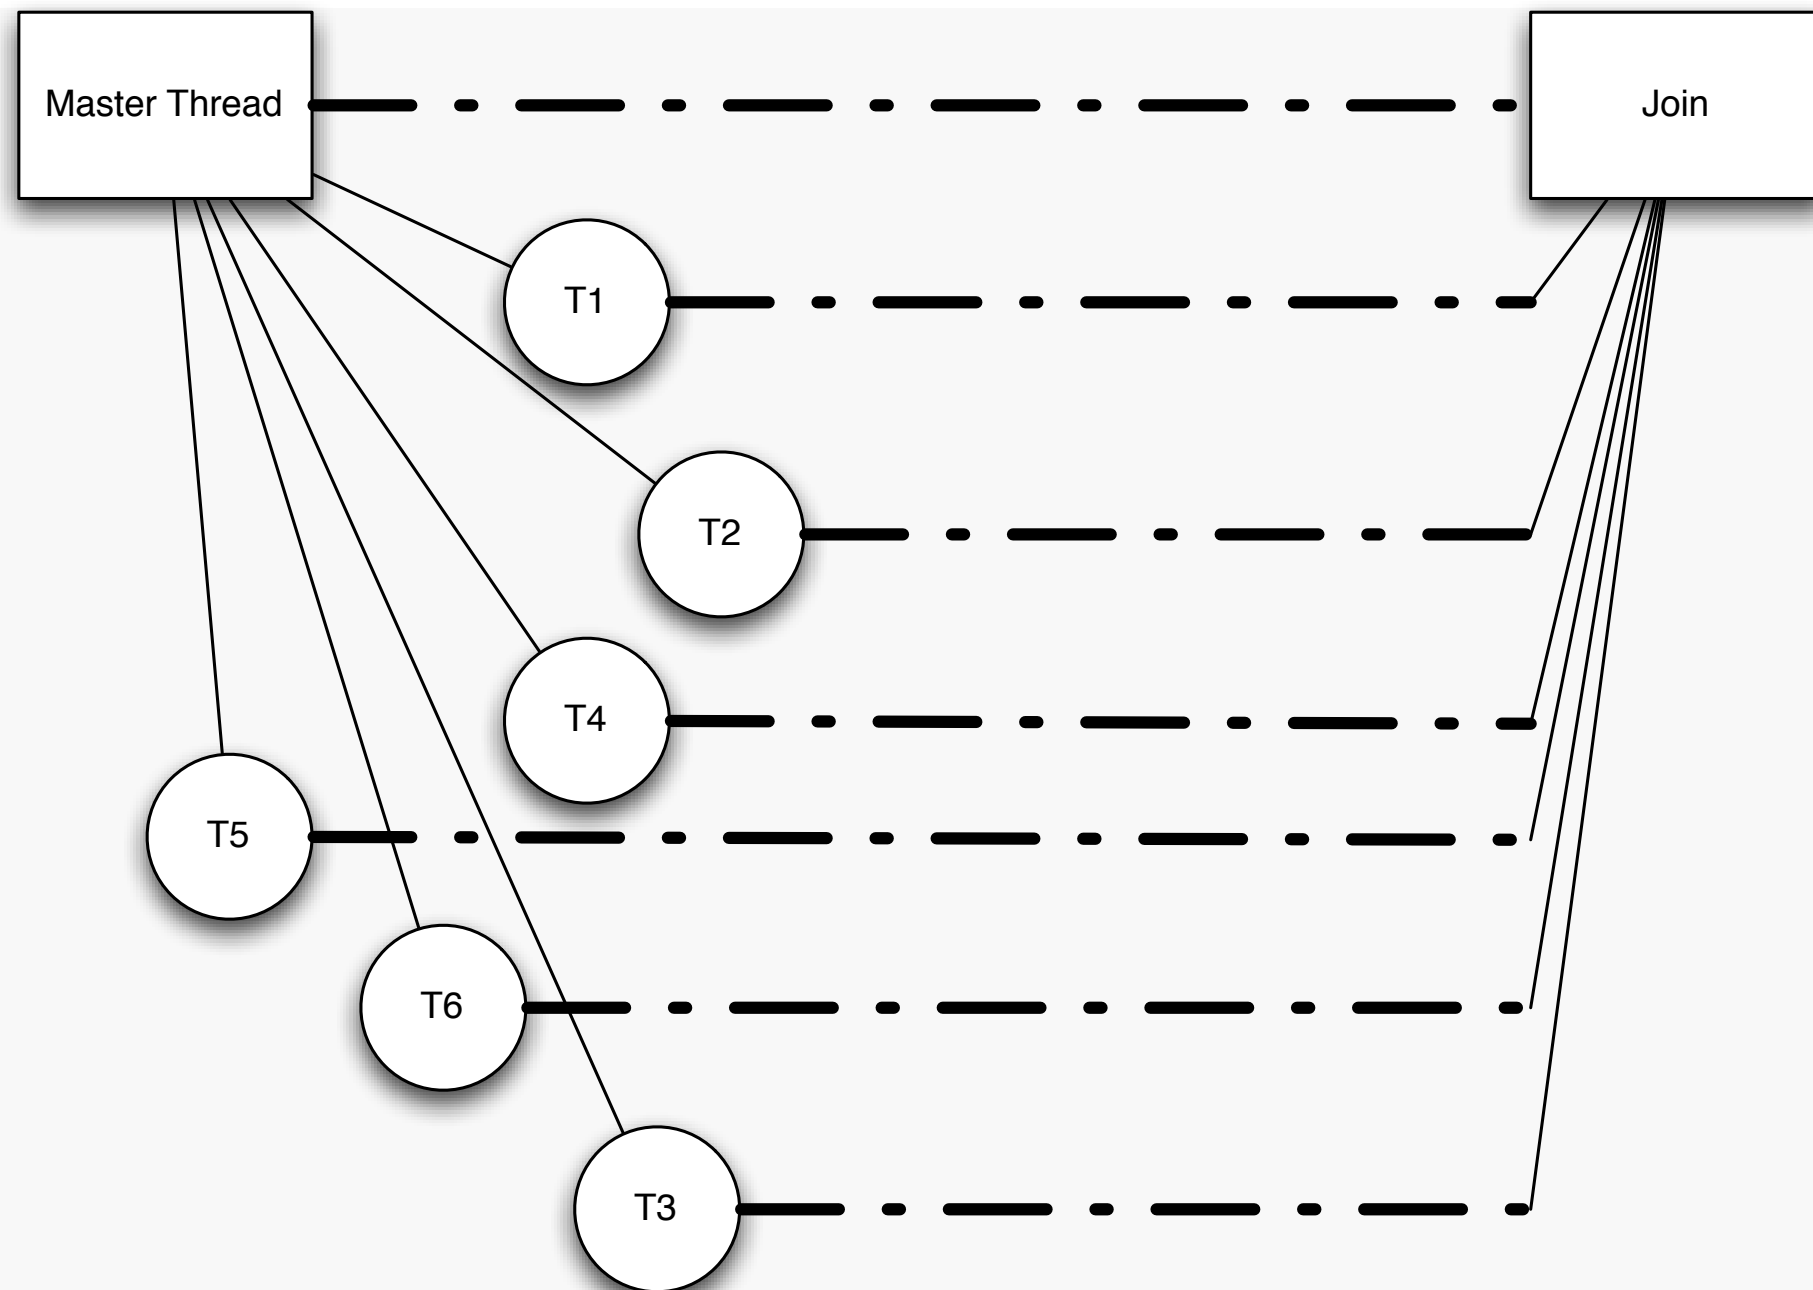


POSIX Threads (2)

- Portable Threading Library across Unix OSes
 - All POSIX-compliant Unixes implement *pthread*
 - Linux, Mac OS X, Solaris, FreeBSD, OpenBSD, etc.
- Also Windows:
 - E.g. Open Source: *pthread-win32*
- BTW: Windows Threads are different!



Six Separate Threads



Creating a pthread

- `#include <pthread.h>`

- `int pthread_create(`

- `pthread_t *thread, // the returned thread id`

- `const pthread_attr_t *attr, // starting attributes`

- `void *(*start_routine)(void*),`

- `// the function to run in the thread`

- `void *arg); // parameter for function`



Where...

- *'thread'* is the ID of the thread
- *'attr'* is the input-only attributes (NULL for standard attributes)
- *'start_routine'* (can be any name) is the function that runs when the thread is started, and which must have the signature:

```
void* *start_routine (void* arg);
```

- *'arg'* is the parameter that is sent to the start routine.
- returns a status code. *'0'* is good, *'-1'* is bad.



Wait for a thread to finish

- `int pthread_join(pthread_t thread, void **value_ptr);`
- where
- '*thread*' is the id of the thread you wish to wait on
- '*value_ptr*' is where the thread's exit status will be placed on exit (NULL if you're not interested.)
- •NB: a thread can be joined only to one other thread!



Hello World -- Creating Threads

```
int main (int argc, const char * argv[]) {  
    pthread_t threads[NUM_THREADS];  
    int rc,t;  
    for (t=0;t<NUM_THREADS;t++) {  
        printf("Creating thread %d\n",t);  
        rc = pthread_create(&threads[t],NULL,  
                           PrintHello,(void *)t);  
        if (rc) {  
            printf("ERROR return code from pthread_create(): %d\n",rc);  
            exit(-1);  
        }  
    }  
}
```



Hello World -- Waiting for Exit

```
// wait for threads to exit
for(t=0;t<NUM_THREADS;t++) {
    pthread_join( threads[t], NULL);
}
```



Thread Code

```
void *PrintHello(void *threadid) {  
    printf("\n%d: Hello World!\n", threadid);  
    pthread_exit(NULL);  
}
```



HelloWorld -- Complete

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS      6

void *PrintHello(void *threadid) {
    printf("\n%d: Hello World!\n", threadid);
    pthread_exit(NULL);
}

int main (int argc, const char * argv[]) {
    pthread_t threads[NUM_THREADS];
    int rc,t;
    for (t=0;t<NUM_THREADS;t++) {
        printf("Creating thread %d\n",t);
        rc = pthread_create(&threads[t],NULL,
                           PrintHello,(void *)t);

        if (rc) {
            printf("ERROR return code from pthread_create(): %d\n",rc);
            exit(-1);
        }
    }

    // wait for threads to exit
    for(t=0;t<NUM_THREADS;t++) {
        pthread_join( threads[t], NULL);
    }
    return(0);
}
```



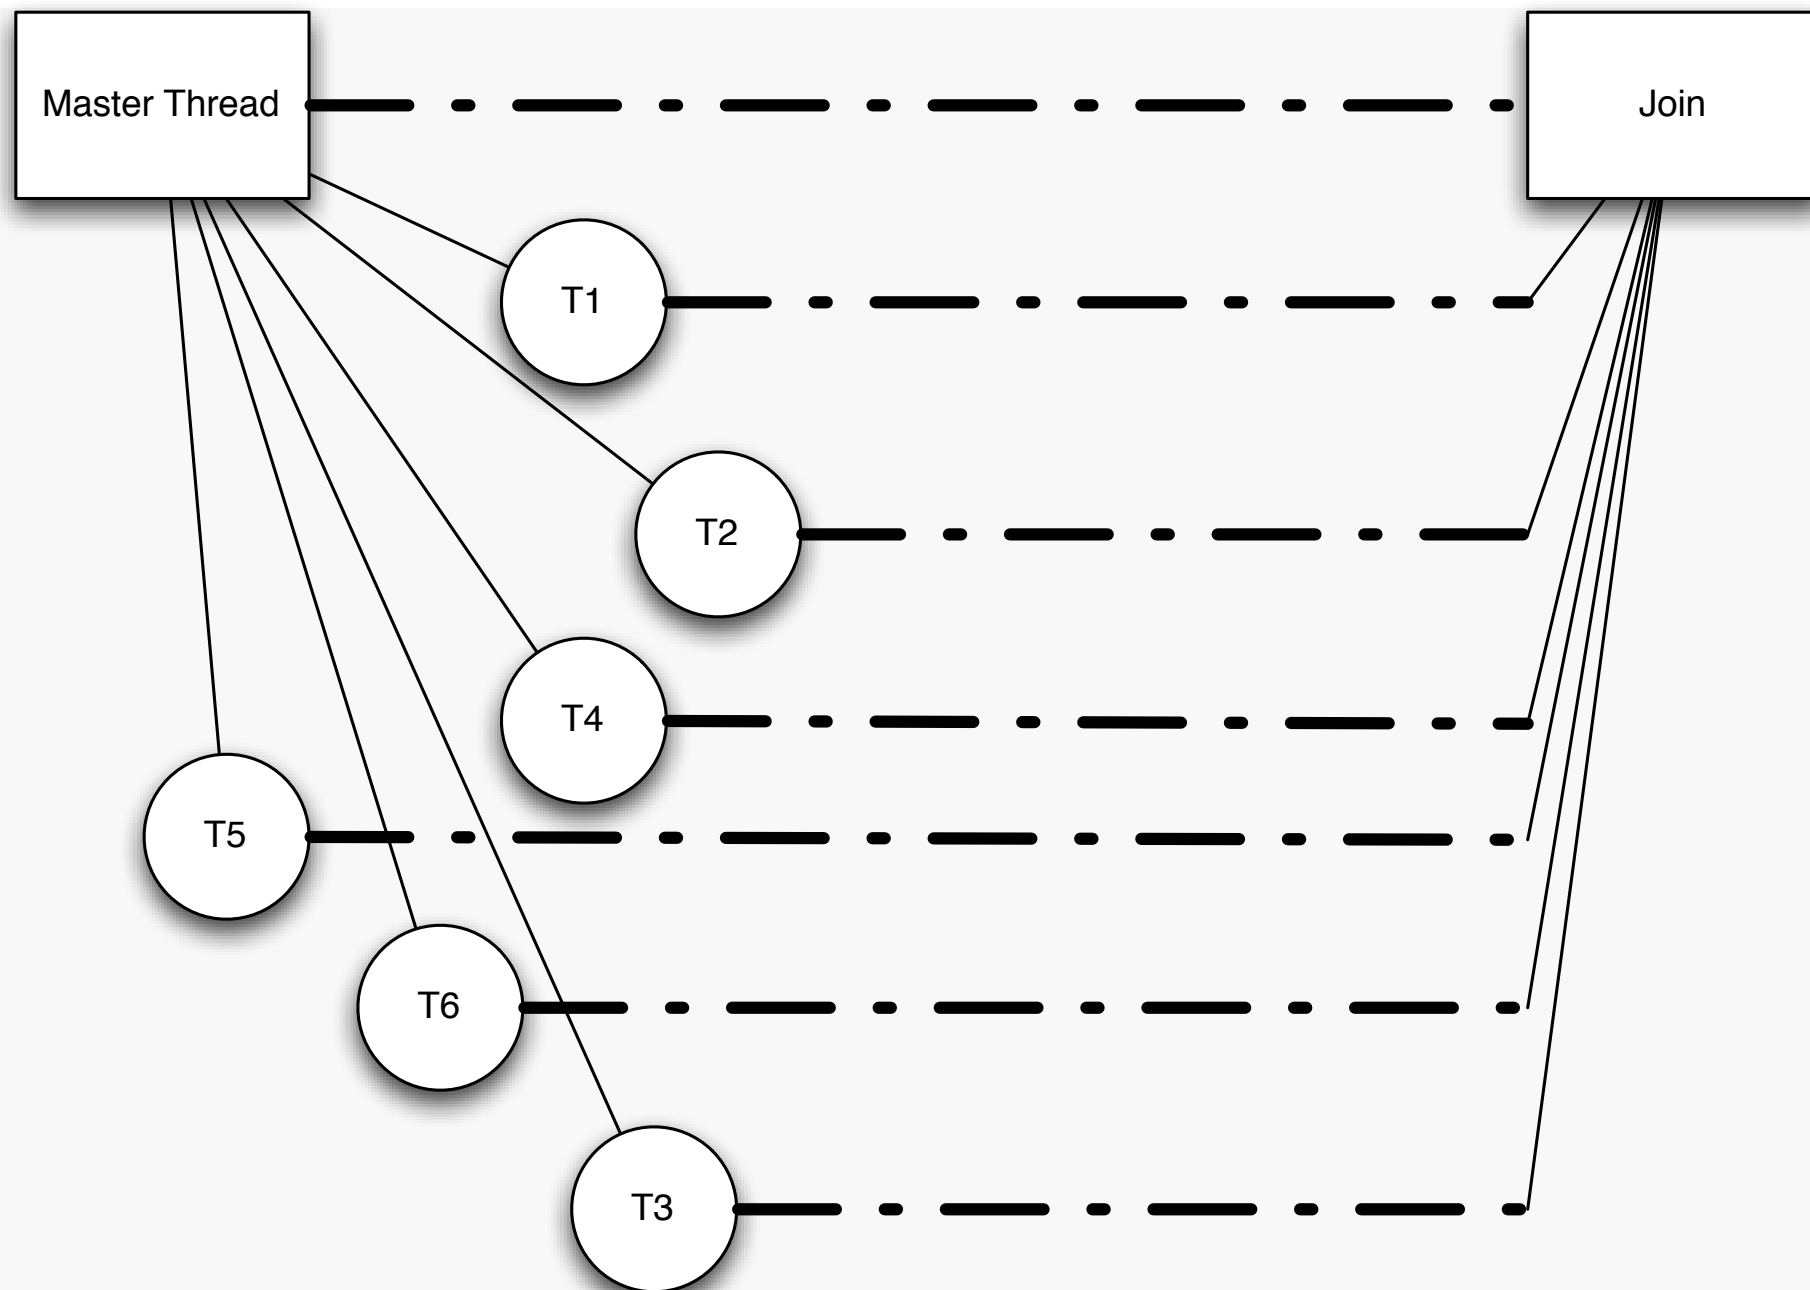
Turing/Stoker/Ubuntu

```
cc -o hello hello.c -pthread
```

- Include the pthread library to allow it to compile and link
 - Sometimes, -lpthread works, but -pthread is correct
- pthread library automatically included in Mac OS X build
- LIVE DEMO



What's curious about this?



What if ..?

- The threads interacted in some way?
 - e.g. each thread increments a global variable that tracks the number of threads that ran?
 - What could possibly go wrong?
- LIVE DEMO



Runtime Behaviour

- The Runtime Behaviour of the Program is no longer under the control of the program.
 - The order in which work gets done on the machine is not exactly under the control of the program
 - It seems to be a price that's paid for parallelism, but
 - What errors can it introduce?
 - Can we prevent them / protect against them / design them out?



CS2016/3D4 plan

- POSIX Thread Programming
 - to expose you to the challenges and problems
- Understanding Concurrency
 - concepts, techniques, modelling, analysis
- Operating Systems (after Study/Review week)

