

Concurrent Systems Operating Systems

3D4 ← → CS2016

Andrew Butterfield
ORI.G39, Andrew.Butterfield@scss.tcd.ie



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

with thanks to Mike Brady

‘Shells’

- A ‘shell’ is a program that allows you to give commands to a computer and get responses.
 - Common GUI-based shells include ‘Finder’ for Mac OS X and ‘Explorer’ for Windows. These are easy to learn but hard to automate.
 - Common UNIX shells are BASH, CSH, TCSH, ASH. These are command-line text-based programs that are a bit more difficult to learn but very powerful indeed. We’ll use BASH, a very-widely used shell in Linux, the BSDs and Mac OS X.
 - The Shell Scripting Primer — search “shell scripting primer pdf”



Important shell principles

- Shell commands generally invoke programs to do the work:
 - e.g. when you write `ls` on the command line, the shell looks for a program called `ls` in a few designated places. Once it finds the program, it executes it.
- Most programs work with standard character-based input sources and output sinks — ‘standard input’, ‘standard output’ and ‘standard error’. These are ‘pipes’ and can be connected to files, the terminal window, other programs, network feeds, etc.
- Most of these programs do something simple but very well — to do complex things, you string programs together.



Combining commands

- Multiple commands:
 - Issue a sequence of commands on one line by separating them with a semicolon (“;”).
- Piping:
 - Using the bar symbol (“|”), you can direct (“pipe”) the standard output of one command into the standard input of the next one.
- Automation:
 - You can write a text file containing sequence of commands and shell commands and constructs. This can be executed as a *shell script*.



Communication

- A parallel program consists of two or more separate threads of execution, that run independently except when they have to interact
- To interact, they must *communicate*
- Communication is typically implemented by
 - sharing memory
 - One thread writes data to memory; the other reads it
 - passing messages
 - One thread sends a message; the other gets it

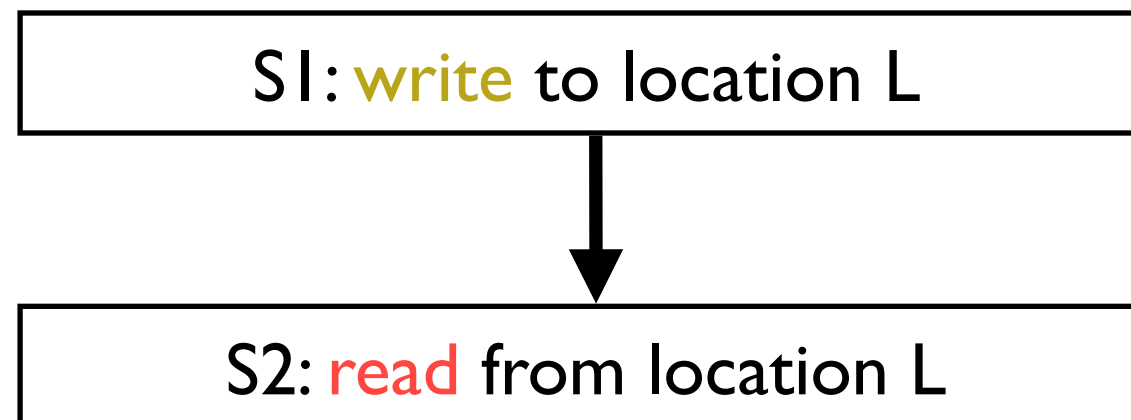


Dependency

- Independent sections of code can run independently.
- We can analyse code for dependencies.
 - To preserve the meaning of the program;
 - To transform the program to reduce dependencies and improve parallelism.



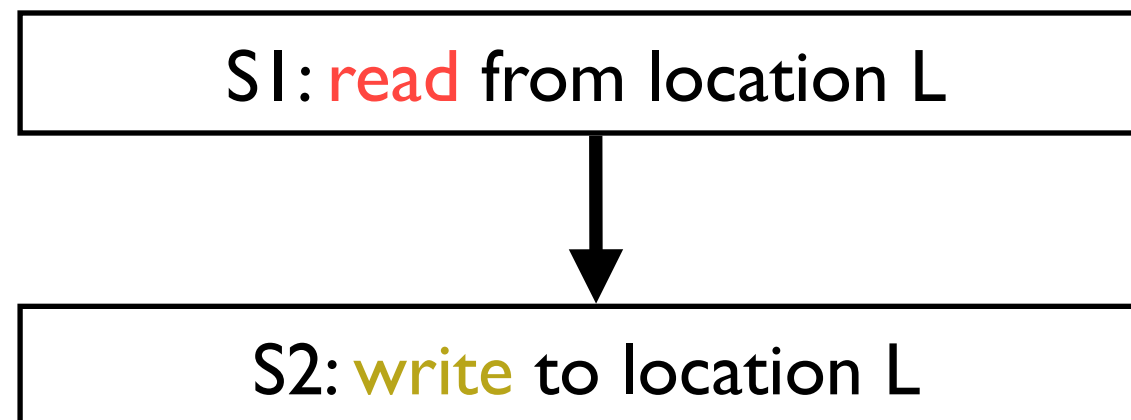
I – Flow Dependence



- Flow Dependence: S2 is flow dependent on S1 because S2 reads a location S1 writes to.
 - It must be written to (S1) before it's read from (S2)
 - Also known as True Dependence, and/or Read-After-Write (RAW)



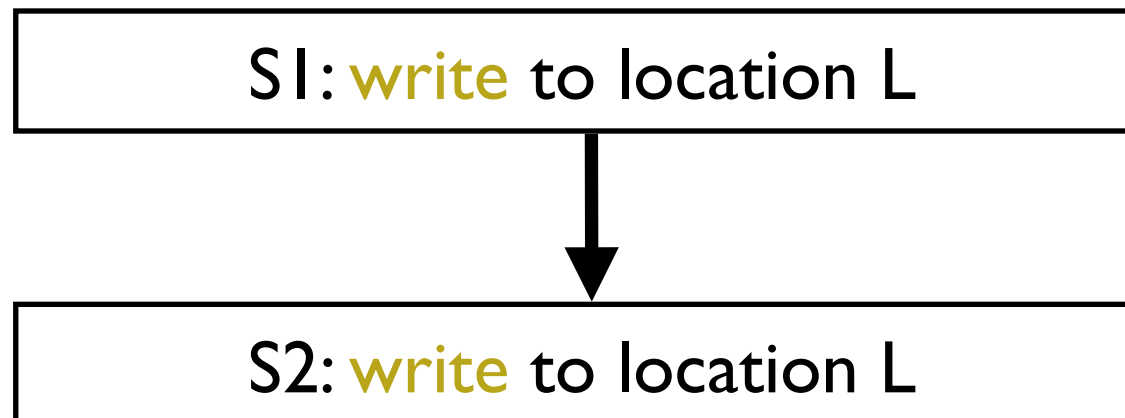
2 – Antidependence



- Antidependence: S2 is antidependent on S1 because S2 writes to a location S1 reads from.
 - It must be read from (S1) before it can be written to (S2)
 - Also known as Write-After-Read (WAR)



3 – Output Dependence




- Output dependence: S2 is output dependent on S1 because S2 writes to a location S1 writes to.
 - The value of L is affected by the order of S1 and S2.
 - Also known as Write-After-Write (WAW)



Example

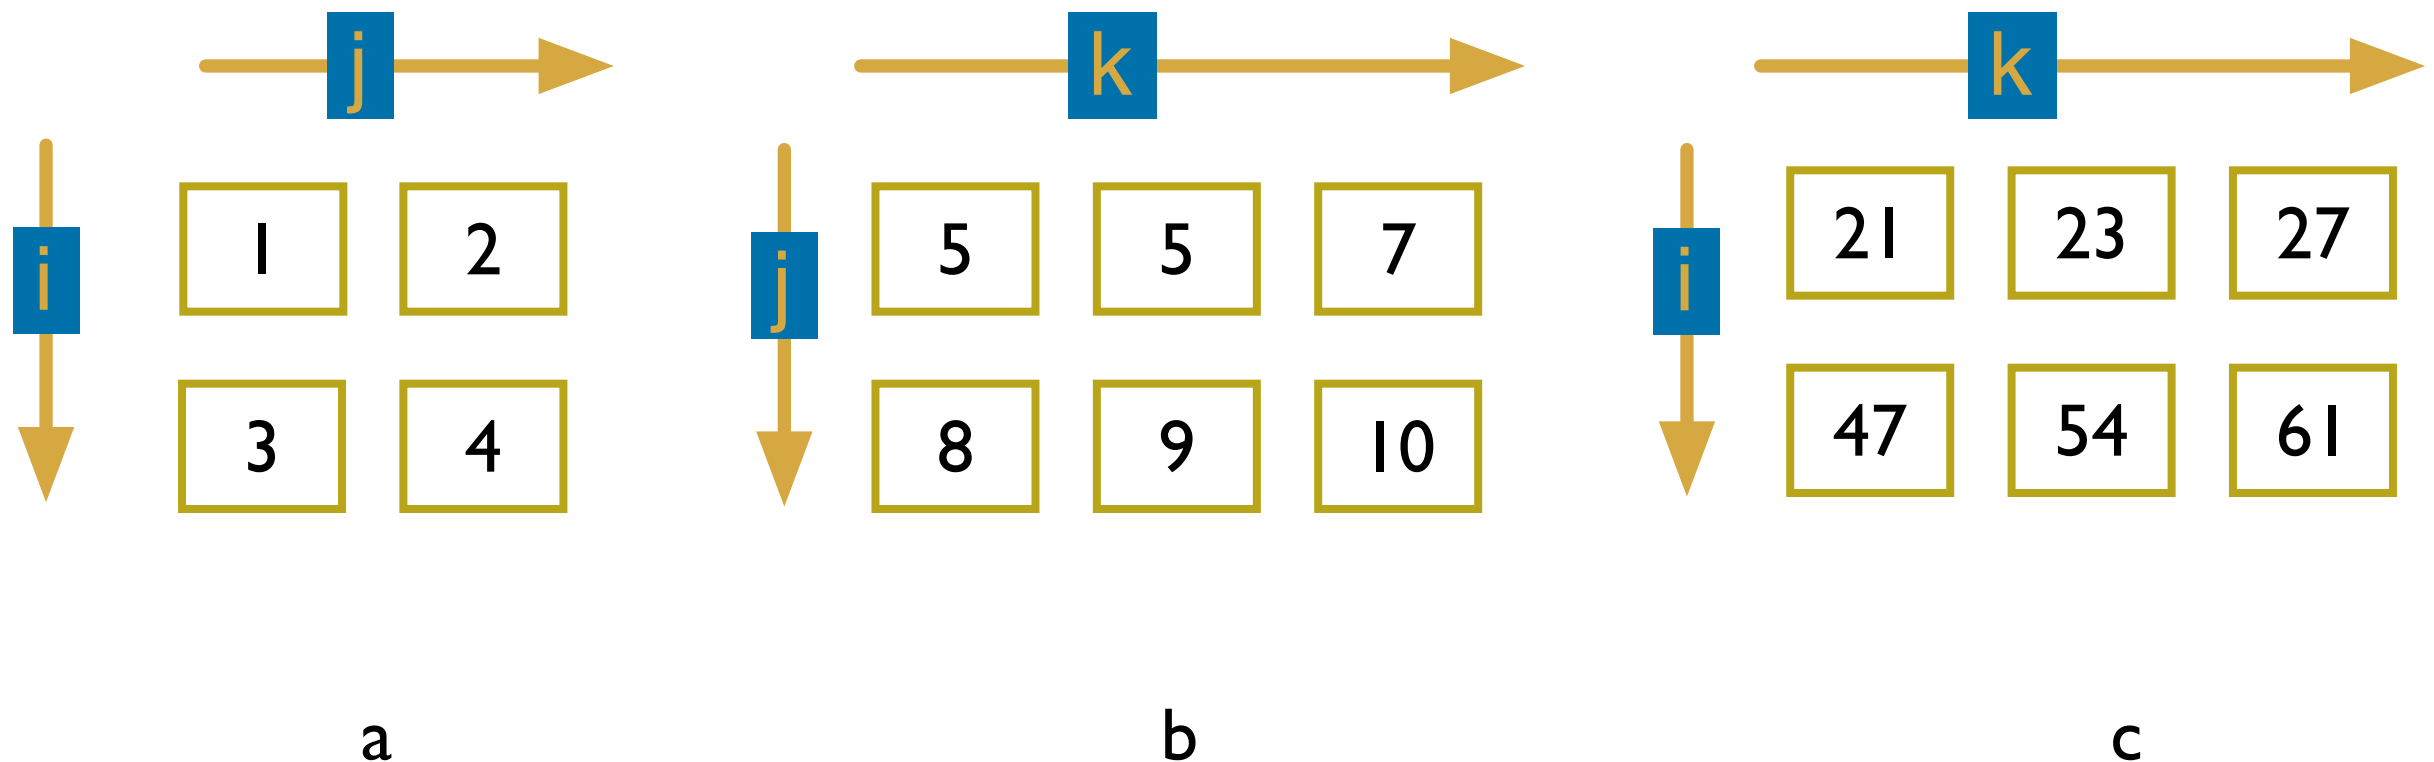
→ Flow
→ Anti
→ Output



Removing Unneeded Dependencies in Loops

- Matrix Multiply. We know, intuitively, we can parallelise this a great deal:

- each element is c is independent



Sample Serial Solution

```
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```



How can we transform it?

- We can work out how to transform it by locating the dependencies in it.
- Some dependencies are intrinsic to the solution, but
- Some are *artefacts* of the *way* we are *solving* the problem;
 - If we can identify them, perhaps we can modify or remove them.



Try three execution agents:

```
with k = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```

```
with k = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```

```
with k = 3;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```



Issues:

- The variable **sum**, as written, is common to all three programs.
- Solution:
 - Make **sum** private to each program to avoid this dependency.



Try Six Execution Agents

```
with k = 1, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 2, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 3, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 1, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 2, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 3, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```



Summarising:

- We could parallelise the original algorithm with some care:
 - Private Variables to avoid unnecessary dependencies
- Actually, we could break this ‘Embarrassingly Parallel’ problem into tiny separate pieces; maybe too small. (*How will we know?*)

