

Concurrent Systems Operating Systems

3D4 ← → CS2016

Andrew Butterfield
ORI.G39, Andrew.Butterfield@scss.tcd.ie



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

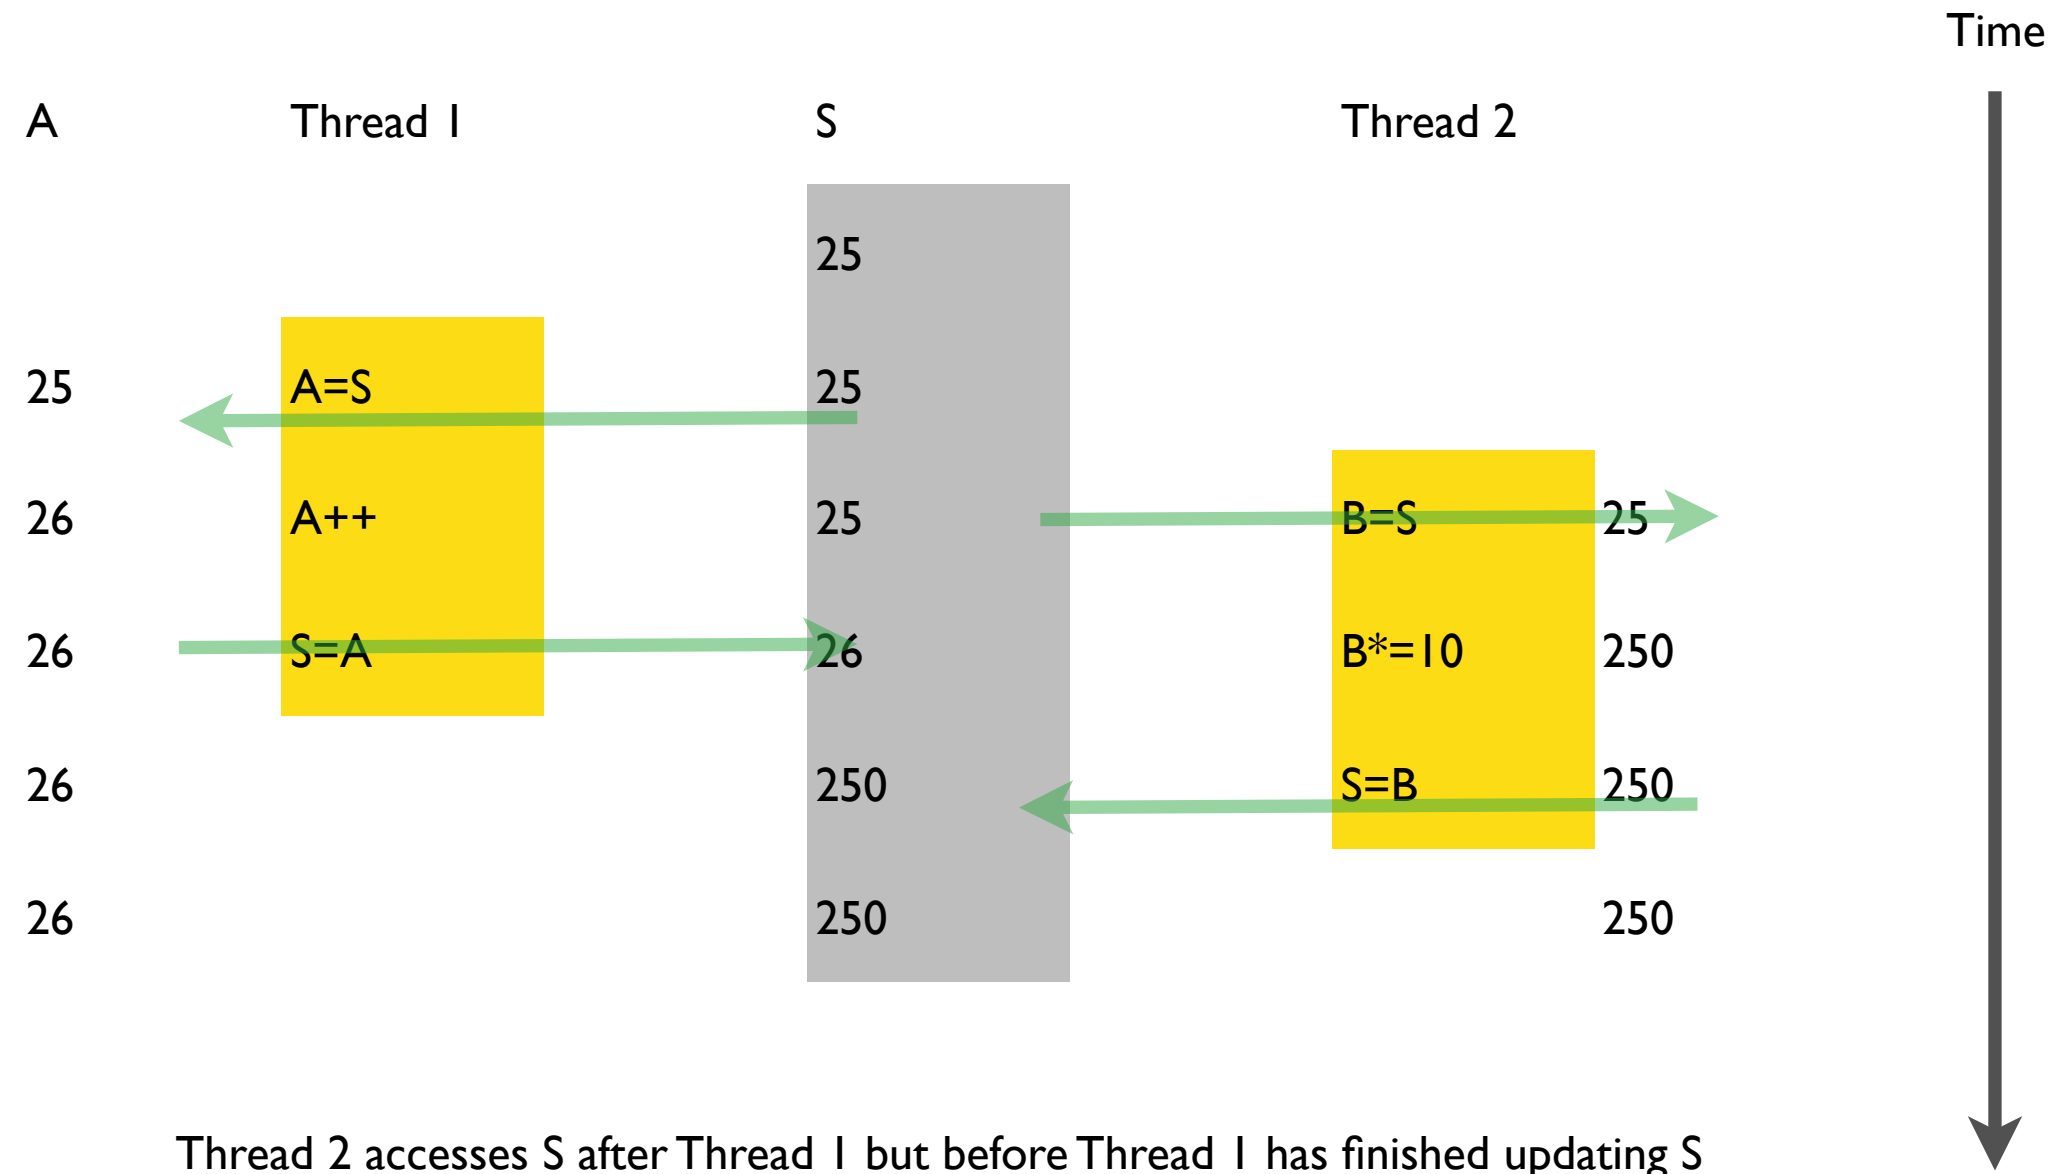
with thanks to Mike Brady

Synchronisation

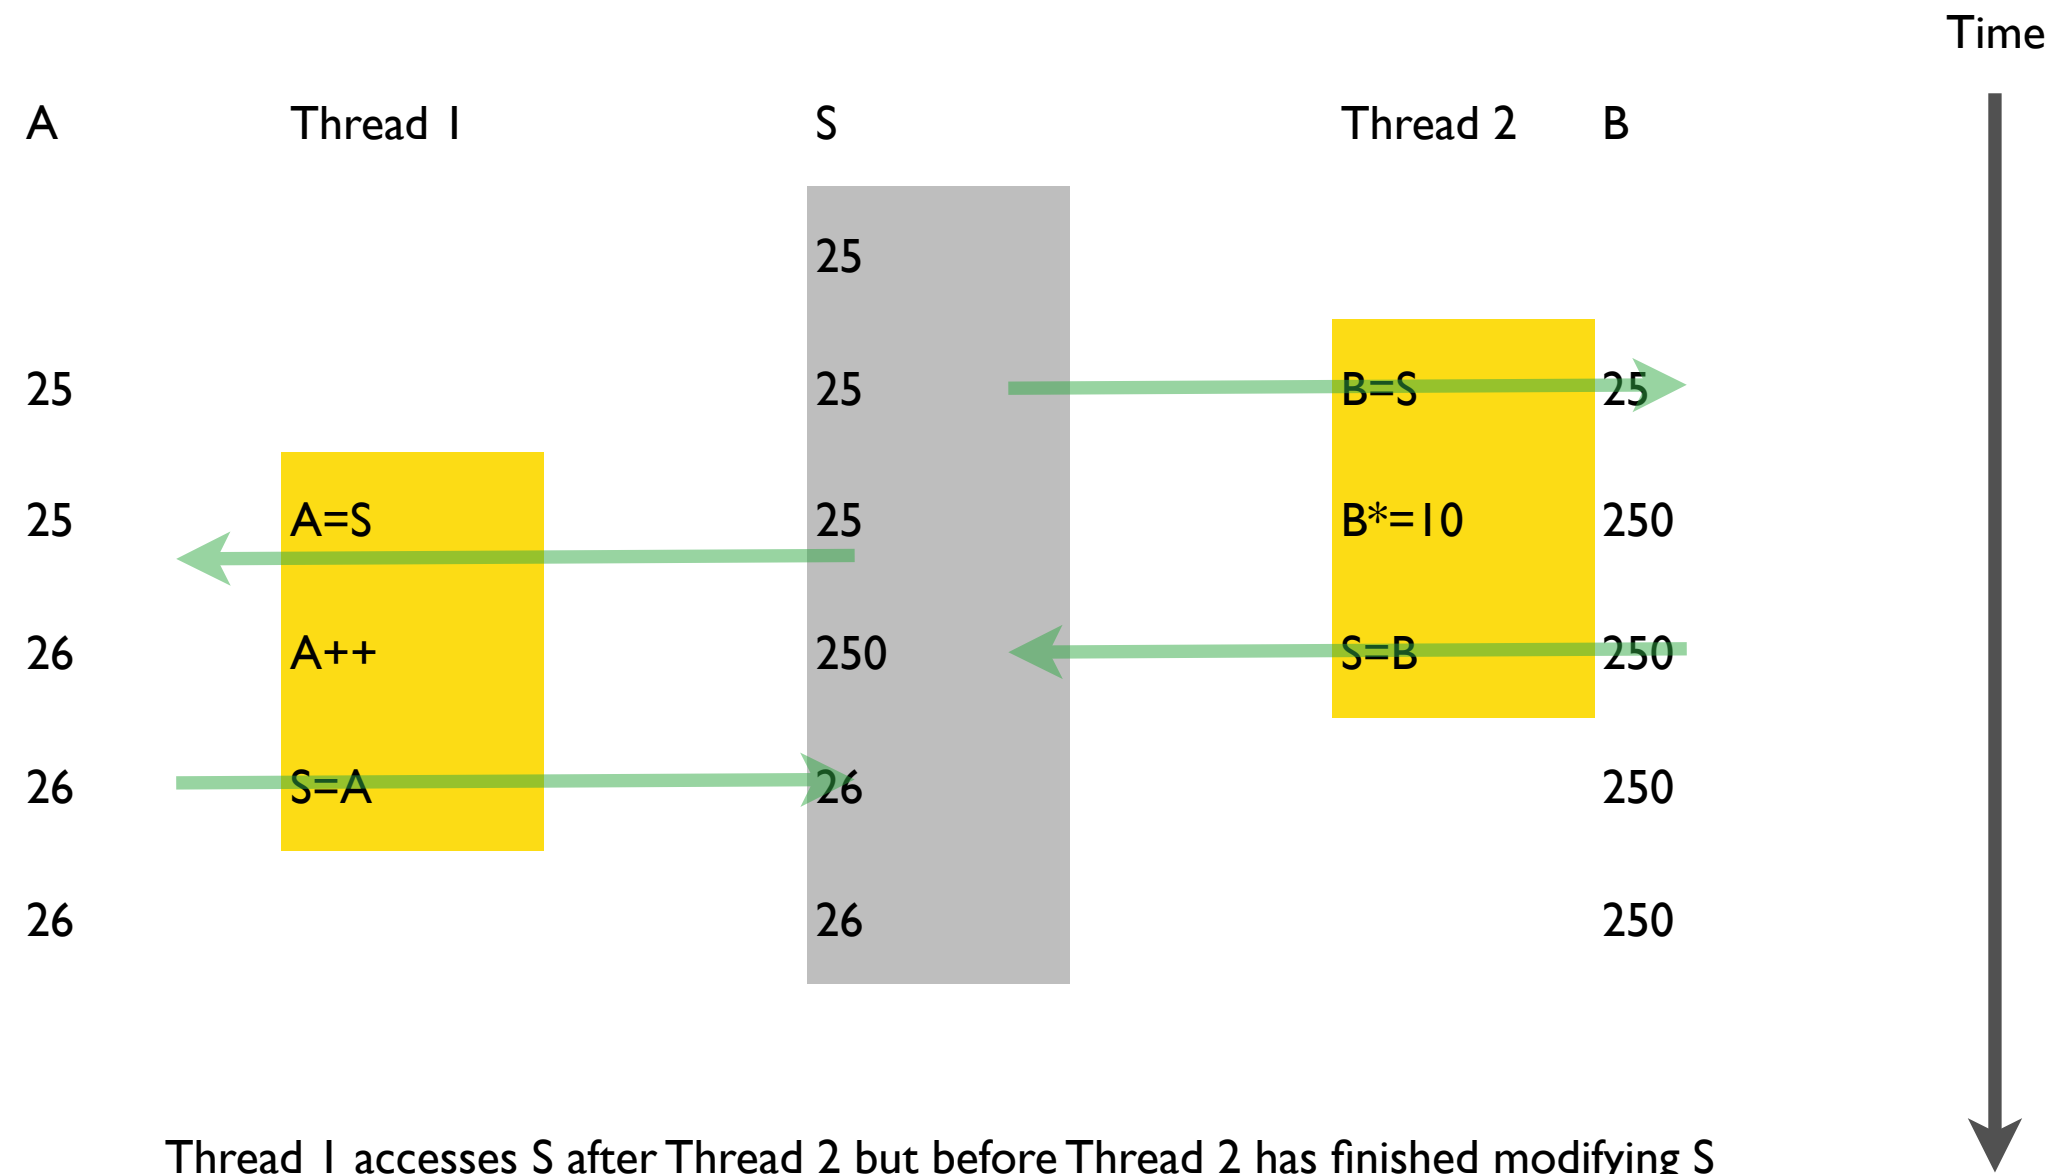
- We've looked at situations where the threads can operate independently -- the 'write sets' of the threads don't intersect.
- Where the write sets intersect, we must ensure that independent thread writes do not damage the data.



Shared Variable S: Thread 1 *before* Thread 2



Shared Variable S: Thread 1 *after* Thread 2



Synchronisation

- Problem: Access to shared resources can be dangerous.
 - These are so-called ‘critical’ accesses.
- Solution. Critical accesses should be made exclusive. Thus, all critical accesses to a resource are *mutually exclusive*.
- In the example, both threads should have asked for exclusive access before making their updates.
 - Depending on timing, one or the other would get exclusive access first. The other would have to wait to get any kind of access.



Mutual Exclusion in pthreads.

- Mutual Exclusion is accomplished using '*mutex*' variables.
- Mutex variables are used to protect access to a resource.

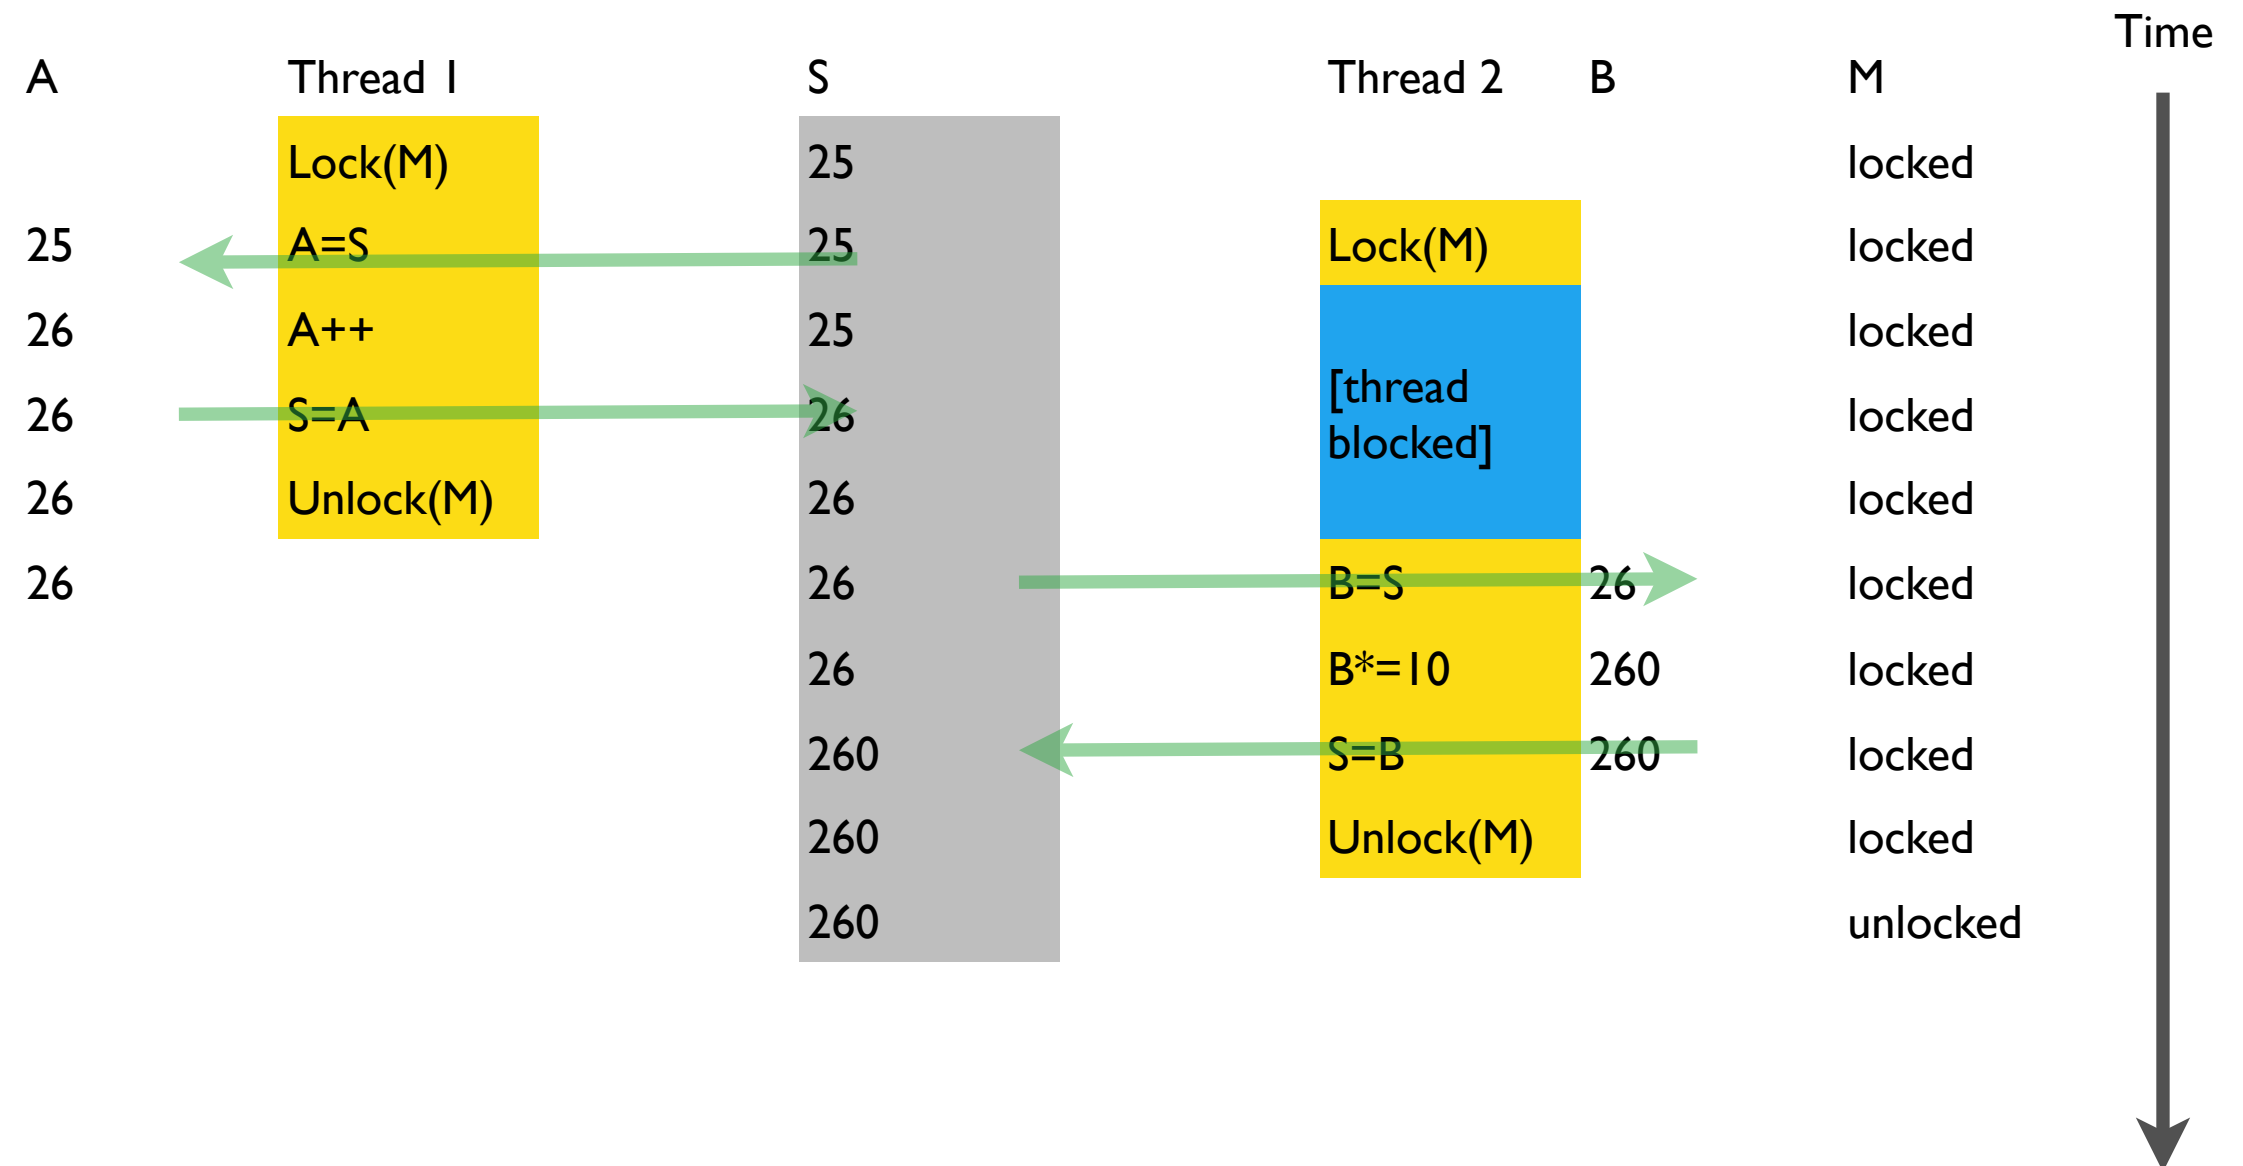


Accessing a protected resource

- To access a mutex-protected resource, an agent must acquire a lock on the mutex variable.
 - If the mutex variable is unlocked, it is immediately locked and the agent has acquired it. When finished, the agent must unlock it.
 - If the mutex variable is already locked, the agent has failed to acquire the lock -- the protected resource is in exclusive use by someone else.
 - The agent is usually blocked until lock is acquired.
 - A non-blocking version of lock acquisition is available.



Shared Variable S Protected by Mutex M



Create pthread mutex variable

- Static:

- `pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;`

- Initially unlocked

- Dynamic

- `pthread_mutex_init(<ref to mutex variable>, attributes)`



Lock and Unlock Mutex

- `pthread_mutex_lock(<mutex variable reference>);`
 - acquire lock or block while waiting
- `pthread_mutex_trylock(<mutex variable reference>);`
 - non-blocking; check returned code
- `pthread_mutex_unlock(<mutex variable reference>);`



Example (1 of 2) – Thread

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define checkResults(string, val) { \
    if (val) { \
        printf("Failed with %d at %s", val, string); \
        exit(1); \
    } \
}

// From http://publib.boulder.ibm.com/infocenter/iserics/v5r3/index.jsp?topic=/rzahw/rzahwe18rx.htm
// Fixed to avoid calls to non-standard pthread_getthreadid_np()

#define NUMTHREADS 3
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
int sharedData=0;
int sharedData2=0;

void *theThread(void *threadid)
{
    int rc;
    printf("Thread %.8x: Entered\n", (int)threadid);
    rc = pthread_mutex_lock(&mutex);
    checkResults("pthread_mutex_lock()\n", rc);

    /***** Critical Section *****/
    printf("Thread %.8x: Start critical section, holding lock\n", (int)threadid);
    /* Access to shared data goes here */
    ++sharedData; --sharedData2;
    printf("Thread %.8x: End critical section, release lock\n", (int)threadid);
    /***** Critical Section *****/

    rc = pthread_mutex_unlock(&mutex);
    checkResults("pthread_mutex_unlock()\n", rc);
    return NULL;
}
```



Example (2 of 2) – Main

```
int main(int argc, char **argv)
{
    pthread_t      thread[NUMTHREADS];
    int            rc=0;
    int            i;

    printf("Enter Testcase - %s\n", argv[0]);

    printf("Hold Mutex to prevent access to shared data\n");
    rc = pthread_mutex_lock(&mutex);
    checkResults("pthread_mutex_lock()\n", rc);

    printf("Create/start threads\n");
    for (i=0; i<NUMTHREADS; ++i) {
        rc = pthread_create(&thread[i], NULL, theThread, (void *)i);
        checkResults("pthread_create()\n", rc);
    }

    printf("Wait a bit until we are 'done' with the shared data\n");
    sleep(3);
    printf("Unlock shared data\n");
    rc = pthread_mutex_unlock(&mutex);
    checkResults("pthread_mutex_lock()\n",rc);

    printf("Wait for the threads to complete, and release their resources\n");
    for (i=0; i <NUMTHREADS; ++i) {
        rc = pthread_join(thread[i], NULL);
        checkResults("pthread_join()\n", rc);
    }

    printf("Results: sharedData: %d, sharedData2: %d\n",sharedData,sharedData2);

    printf("Clean up\n");
    rc = pthread_mutex_destroy(&mutex);
    printf("Main completed\n");
    return 0;
}
```



Problems

- Voluntary

- Mutexes ‘protect’ code.
- Other programmers don’t have to use them to get access to the variables the code accesses.
 - This is part of the tradeoff. Use processes rather than threads if you want better protection.

- Unfair

- If multiple threads are blocked on a mutex, the order in which they waken up is not guaranteed to be any particular order.

