

# Concurrent Systems Operating Systems

3D4 ← → CS2016

*Andrew Butterfield*  
*ORI.G39, [Andrew.Butterfield@scss.tcd.ie](mailto:Andrew.Butterfield@scss.tcd.ie)*



**Trinity College Dublin**  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

*with thanks to Mike Brady*

# Problem Tutorial

- The *transpose* of a matrix **M** is a matrix **T**, where:  
 $T[i, j] = M[j, i]$  for all  $i, j$ .
- Consider a sequential program to transpose a matrix.
- How would you parallelise it?
- What problems/issues do you foresee?



# Running Sequential Code

- Ran sequential code for a large matrix, with and without `printf` statements
- Source code: `seq-transpose-m2t.c`
- Sizing: `#define MSIZE 3000`
- Compiling: `cc -o seq seq-transpose-m2t.c`
- Running (and timing): `time ./seq`



# Running Concurrent Code

- Developed code to use a thread per row.
- Ran concurrent code for a large matrix, with and without `printf` statements
- Source code: `row-transpose-m2t.c`
- Sizing: `#define MSIZE 3000`
- Compiling: `cc -o row row-transpose-m2t.c`
- Running (and timing): `time ./row`



# Timing Comparison

- We have four runs, size=3000,  
sequential and row-concurrent, with and without `printf` output

prog.	printf?	real	user	sys	
seq	yes	24.169	9.459	6.687	
seq	no	0.125	0.106	0.018	
row	yes	52.813	28.246	201.822	4.35
row	no	0.166	0.154	0.283	2.63

4Ghz Intel i7, 32Gb, OS X Mojave

Printing takes a lot of time !

Why is sys time greater than real time for prog `row`?



# Problem Tutorial (revisited)

- Assume that matrix **M** is square
- Consider a program that transposes this matrix ***in place***.
  - $M[i, j] = M[j, i]$  for all  $i, j$ .
- How would you parallelise it now?
- What additional problems/issues do you foresee?



# In-place update discussion

- No code was developed - just a general discussion of issues
- Looks fully independent at first glance, so how about a thread  $t(i,j)$  for each value of  $i$  and  $j$  each doing  $M[i][j] = M[j][i];$  ?
  - $MSIZE*MSIZE$  threads! Too much thread create/memory/join overhead
  - if  $t(i,j)$  runs before  $t(j,i)$  the final outcome is  $M[i,j] = M[j,i] =$  initial value of  $M[j,i]$
- Another solution: threads  $t(i,j)$  where  $i < j$ , that swap  $M[i,j]$  and  $M[j,i]$ 
  - $(MSIZE*(MSIZE-1))/2$  threads!
- More like the **row** program: a thread  $t(i)$  that works for all  $j > i$ ?
  - better, but some threads do 1 or 2 swaps, while others do almost  $MSIZE$  swaps
  - Suggestion:  $t(ia,ib)$  that does both  $t(ia)$  and  $t(ib)$  as above, but  $ia+ib = MSIZE-2$

