

# Concurrent Systems Operating Systems

3D4  CS2016

*Andrew Butterfield*  
*ORI.G39, [Andrew.Butterfield@scss.tcd.ie](mailto:Andrew.Butterfield@scss.tcd.ie)*



**Trinity College Dublin**  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

*with thanks to Mike Brady*

# Linear Temporal Logic

- Logic that reasons about state-change over time
  - Allows us to write properties that apply to an entire scenario, i.e., sequences of states
- LTL is the simplest of a large family of temporal (and "modal") logics
  - Also often used to reason about computing systems are CTL (Computational Tree Logic) and CTL\*
- LTL is an extension of standard propositional ("digital") logic with temporal operators
- We use the SPIN LTL notation here
  - there are more mathematical forms in the literature
  - LTL notation text notation is not standard - Wikipedia has a common but different variant.



# Propositional Logic in LTL

- LTL contains propositional logic with all the usual parts:
  - Logical constants: `true`, `false`
  - Variables: `a`, `b`, .. `p`, `q`, `xx`, `yY`, ... - they must start with lowercase letter
  - Negation: `!`      Logical-and: `&&` (also `/\` )      Logical-or: `||` (also `V` )
  - Logical Implication: `->`
  - Logical Equivalence: `<->`



# #define vs. inline

- SPIN uses the C pre-processor (CPP) to process Promela files
  - So all the CPP facilities are available, such as `#define`, `#if`, `#ifdef`, etc.
- In the producer-consumer example (`prodcons2.pml`), we used something similar called `inline`.
  - What is the difference?
- If an error occurs in code produced by a macro defined using `#define`,
  - the error is reported at the point of use in the expanded macro text
- If an error occurs in code produced by a macro defined using `inline`,
  - the error is reported at the relevant line in the `inline` definition itself
    - Generally much more useful.



# State(s) in LTL

- A state is defined by the values of a given set of variables.
  - propositional expressions evaluate to true or false based on those values in a given *single* state
- LTL in general is interested in sequences (a.k.a. paths) of states.
  - In general, an LTL expression is deemed to be true of a *given starting* state in such a path:
    - if it holds true for the path from that state onwards.
- We will denote such a path as a sequence of indexed state:s
  - $s_0, s_1, s_2, \dots, s_{i-1}, s_i, s_{i+1}, \dots$
  - A temporal property is true "at" state  $s_i$  if it is true for the path starting with  $s_i$ .



# Linear Temporal Operators (Until)

- For most applications, all we need to do is to define one temporal operator ("until")
  - The rest of the operators can be expressed in terms of this
- LTL starts with a notion of "weak until", which says that  $(p \text{ until } q)$  is true at starting state  $s_i$  if
  - $q$  is true at  $s_i$ , or
  - $p$  is true at  $s_i$  and  $(p \text{ until } q)$  is true at  $s_{i+1}$
  - Note that there is no requirement for  $q$  to ever be true, but if it is never so, then  $p$  must always be true
- Strong "until" is weak until with an additional requirement that  $q$  must become true after a finite number of steps.
- SPIN uses the notation  $U$  to denote strong until - it does not have the weak operator.



# Linear Temporal Operators (derived)

- We can now derive two other useful operators
  - one using "weak-until", the other using "strong-until"
- Always  $p \ ([ \ ] p)$  is true if  $p$  is true in every state in the path
  - It can be defined using weak-until as  $[ \ ] p = p \text{ until } \text{false}.$
- Eventually  $p \ ( < > p)$  is true if  $p$  is true at least once, somewhere along path, after a finite number of steps
  - It is defined using strong-until as  $< > p = \text{true} \ U \ p$



# Compiling LTL with SPIN

- SPIN can take a (quoted) LTL formula as an argument on the command-line and output the corresponding never claim on stdout :  
(Here, for example, we are claiming that p and q are never true together at the start)

```
[:- spin -f "p && q"
never {      /* p && q */
accept_init:
T0_init:
    do
        :: atomic { ((p) && (q)) -> assert(!((p) && (q))) }
    od;
accept_all:
    skip
}
```

- We can also pipe the formula from a file:

```
[:- spin -f "$(< pandq.txt) "
```





# Where are the comparisons?

- The LTL notation supported by SPIN does not contain useful comparison operators such as `==`, `<=`, `>=`
- It also doesn't support general expressions like `x+1` or `y * (z - x)`, or even `x+1 < y`.
- Instead, we must use the C-preprocessor `#define` to use a single variable to denote a boolean expression:
  - `#define p (x+1<y) .`  
The parentheses are recommended !
  - We can then use `p` in LTL to stand for `x+1<y`.
  - These `#defines` go in the Promela file into which the `never` claim is put.



# Common LTL Predicates

LTL	Reads as...	Property
$[ ]p$	always p	invariance
$<>p$	eventually p	guarantee
$p \rightarrow <>q$	p implies eventually q	response
$p \rightarrow q \text{ U } r$	p implies q until r	precedence
$[ ]<>p$	always eventually p	recurrence (progress)
$<>[ ]p$	eventually always p	stability (non-progress)
$<>p \rightarrow <>q$	eventually p implies eventually q	correlation

