

Concurrent Systems Operating Systems

3D4 ← → CS2016

Andrew Butterfield
ORI.G39, Andrew.Butterfield@scss.tcd.ie



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

with thanks to Mike Brady

Scheduler Issues

- Fairness
 - Make sure every process gets processor time eventually.
- Efficiency
 - Make best use of the resources available.
- Timeliness
 - Respond quickly to events, e.g. keyboard, etc.
 - Provide processor time reliably, in real-time, over extended periods.
- Requirements may conflict.



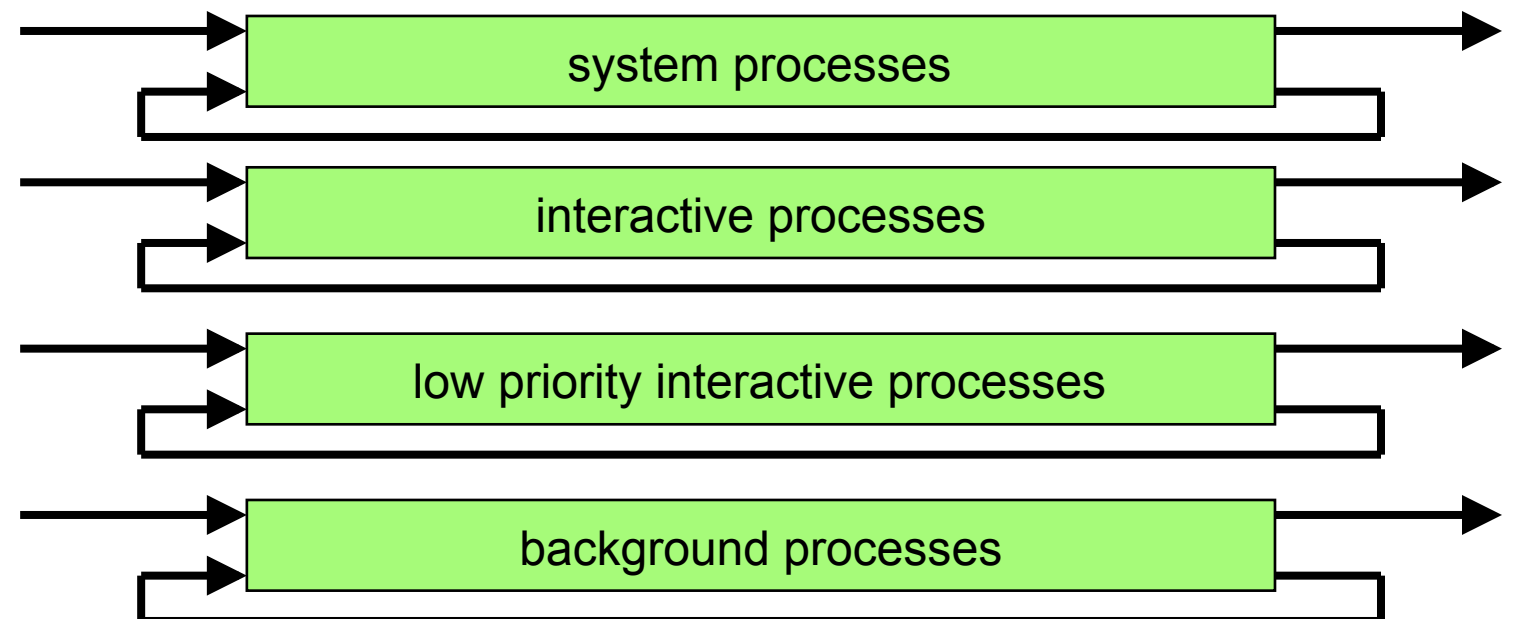
Real-Time

- Real-time means the correctness or utility of a program depends on being logically correct but also upon the time by which results or services are provided by the program. Real-time systems have deadlines.
 - *Hard* real-time. If a deadline is missed, the system fails totally.
 - *Firm* real-time. If a deadline is missed, the delayed service is useless, but the overall system can tolerate the occasional failure.
 - *Soft* real-time. If a deadline is missed, the delayed service is degraded, not useless, and the overall system is degraded but still useful/usable.
- Windows, Mac OS X, [Standard] Linuxes & FreeBSD/OpenBSD/NetBSDs are *not* hard real-time-capable OSes.



Scheduling Arrangements

- Multi-level queue scheduling
 - Often we can group processes together based on their characteristics or purpose
 - system processes
 - interactive process
 - low priority interactive processes
 - background processes



Scheduling Arrangements

- Multi-level queue scheduling
 - Processes are permanently assigned to these groups and each group has an associated queue
 - Each process group queue has an associated queuing discipline
 - FCFS, SJF, priority queuing, Round-Robin, ...
 - Additional scheduling is performed between the queues
 - Often implemented as fixed-priority pre-emptive scheduling
 - Each queue may have absolute priority over lower priority queues
 - In other words, no process in the background queue can execute while there are processes in the interactive queue
 - If a background process is executing when an interactive process arrives in the interactive queue, the background process will be preempted and the interactive process will be scheduled
 - Alternatively, each queue could be assigned a time slice
 - For example, we might allow the interactive queue 80% of the CPU time, leaving the remaining 20% for the other



Scheduling Arrangements

- Multi-level feedback queue scheduling
 - Unlike multi-level queue scheduling, this scheme allows processes to move between queues
 - Allows I/O bound processes to be separated from CPU bound processes for the purposes of scheduling
 - Processes using a lot of CPU time are moved to lower priority queues
 - Processes that are I/O bound remain in the higher priority queues
 - Processes with long waiting times are moved to higher priority queues
 - Aging processes in this way prevents starvation
 - Higher priority queues have absolute priority over lower priority queues
 - A process arriving in a high-priority queue will preempt a running process in a lower priority queue




Scheduling Arrangements

- Multi-level feedback queue scheduling

- When a process arrives on the ready queue (a new process or a process that has completed waiting for an event), it is placed on the highest priority queue
- Processes in the highest priority queue have a quantum of 8ms
- If a high priority process does not end its CPU burst within 8ms, it is preempted and placed on the next lowest priority queue
- When the highest priority queue is empty, processes in the next queue are given 16ms time slices, and the same behaviour applies
- Finally, processes in the bottom queue are serviced in FCFS order

- We can change

- number of queues
- queuing discipline
- rules for moving process to higher or lower priority queues



Multi-processor scheduling

- Flynn's taxonomy of parallel processor systems
 - Single instruction single data (SISD)
 - Single instruction multiple data (SIMD) - e.g., (some) “graphics cards”
 - vector / array processor
 - Multiple instruction single data (MISD)
 - argument over precise meaning (pipelined processor ???)
 - Multiple instruction multiple data (MIMD) - e.g. “multicore”
 - shared memory (tightly coupled)
 - asymmetric (master/slave)
 - symmetric (SMP)
 - distributed memory (loosely coupled)
 - cluster



Multi-processor scheduling

- Processor affinity
 - If a process can execute on any processor, what are the implications for cache performance?
 - Most SMP systems implement a processor affinity scheme
 - soft affinity: an attempt is made to schedule a process to execute on the CPU it last executed on
 - hard affinity: processes are always executed on the same CPU



Multi-processor scheduling

- Load balancing
 - Scheduler can use a common ready queue for all processors or use private per-processor ready queues
 - If per-processor ready queues are used (and no affinity or soft-affinity is used), then load-balancing between CPUs becomes an issue
 - Two approaches:
 - pull-migration: idle processors pull ready jobs from other processors
 - push-migration: load on each processor is monitored and jobs are pushed from overloaded processors to under-loaded processors
- Load-balancing competes with affinity



Scheduling in Linux (2.6...)

- Three categories of process:
 - Interactive
 - Batch
 - Real-time



Process scheduling in Linux

- Scheduling goals
 - Fast response time
 - High throughput
 - Avoid starvation
 - Satisfy the needs of both high- and low-priority processes
 - ...



Process scheduling in Linux

- Pre-emptive scheduling
 - When quantum of the executing process expires, another process may be executed in its place



Process scheduling in Linux

- Processes are ranked according to priorities
 - Priorities are dynamic: the operating system scheduler adjusts process priorities based on behaviour (e.g. decrease priority of long-running processes, boost priority when on I/O completion, etc.)
 - This results in complex schemes for determining the current priority for a process but simpler schemes for selecting the next process to execute.



Process scheduling in Linux


- Every (conventional) process has a static priority
 - in the range 100 (highest priority) ... 139 (lowest priority)
- Static priority determines base time quantum

if (static priority < 120) then
 base time quantum = (140 – static priority) x 20
else
 base time quantum = (140 – static priority) x 5



Process scheduling in Linux

base time quantum (ms)



Process scheduling in Linux

- Dynamic priority determines which process to execute:
 - $\text{dynamic priority} = \max(100, \min(\text{static priority} - \text{bonus} + 5, 139))$
 - bonus in the range 0 ... 10 but the effect of a “bonus” less than 5 is to lower the dynamic priority of the process
 - bonus is based on average sleep time and is generated from a lookup table that gives a bigger bonus to processes that sleep for longer



Process scheduling in Linux

Average sleep time	Bonus
$0 \leq \text{average sleep time} < 100$	0
$100 \leq \text{average sleep time} < 200$	1
$200 \leq \text{average sleep time} < 300$	2
$300 \leq \text{average sleep time} < 400$	3
$400 \leq \text{average sleep time} < 500$	4
$500 \leq \text{average sleep time} < 600$	5
$600 \leq \text{average sleep time} < 700$	6
$700 \leq \text{average sleep time} < 800$	7
$800 \leq \text{average sleep time} < 900$	8
$900 \leq \text{average sleep time} < 1000$	9
$1000 \leq \text{average sleep time}$	10



Process scheduling in Linux

- Linux differentiates between interactive and batch processes
- A process is interactive if $(\text{bonus} - 5 \geq \text{static priority} / 4 - 28)$
 - It is easier for a high priority process to become interactive
 - Very low priority processes cannot become interactive
- To avoid starvation of low-priority process by high-priority processes, all runnable processes are classified as **active** or **expired**
 - **active** processes have not exhausted their quantum and are allowed to run
 - **expired** processes have exhausted their quantum and are not allowed to run until there are no more active processes
 - in practice the situation is a little more complex (e.g. interactive processes can get placed back on the active list as long as there are no “starving” expired processes)

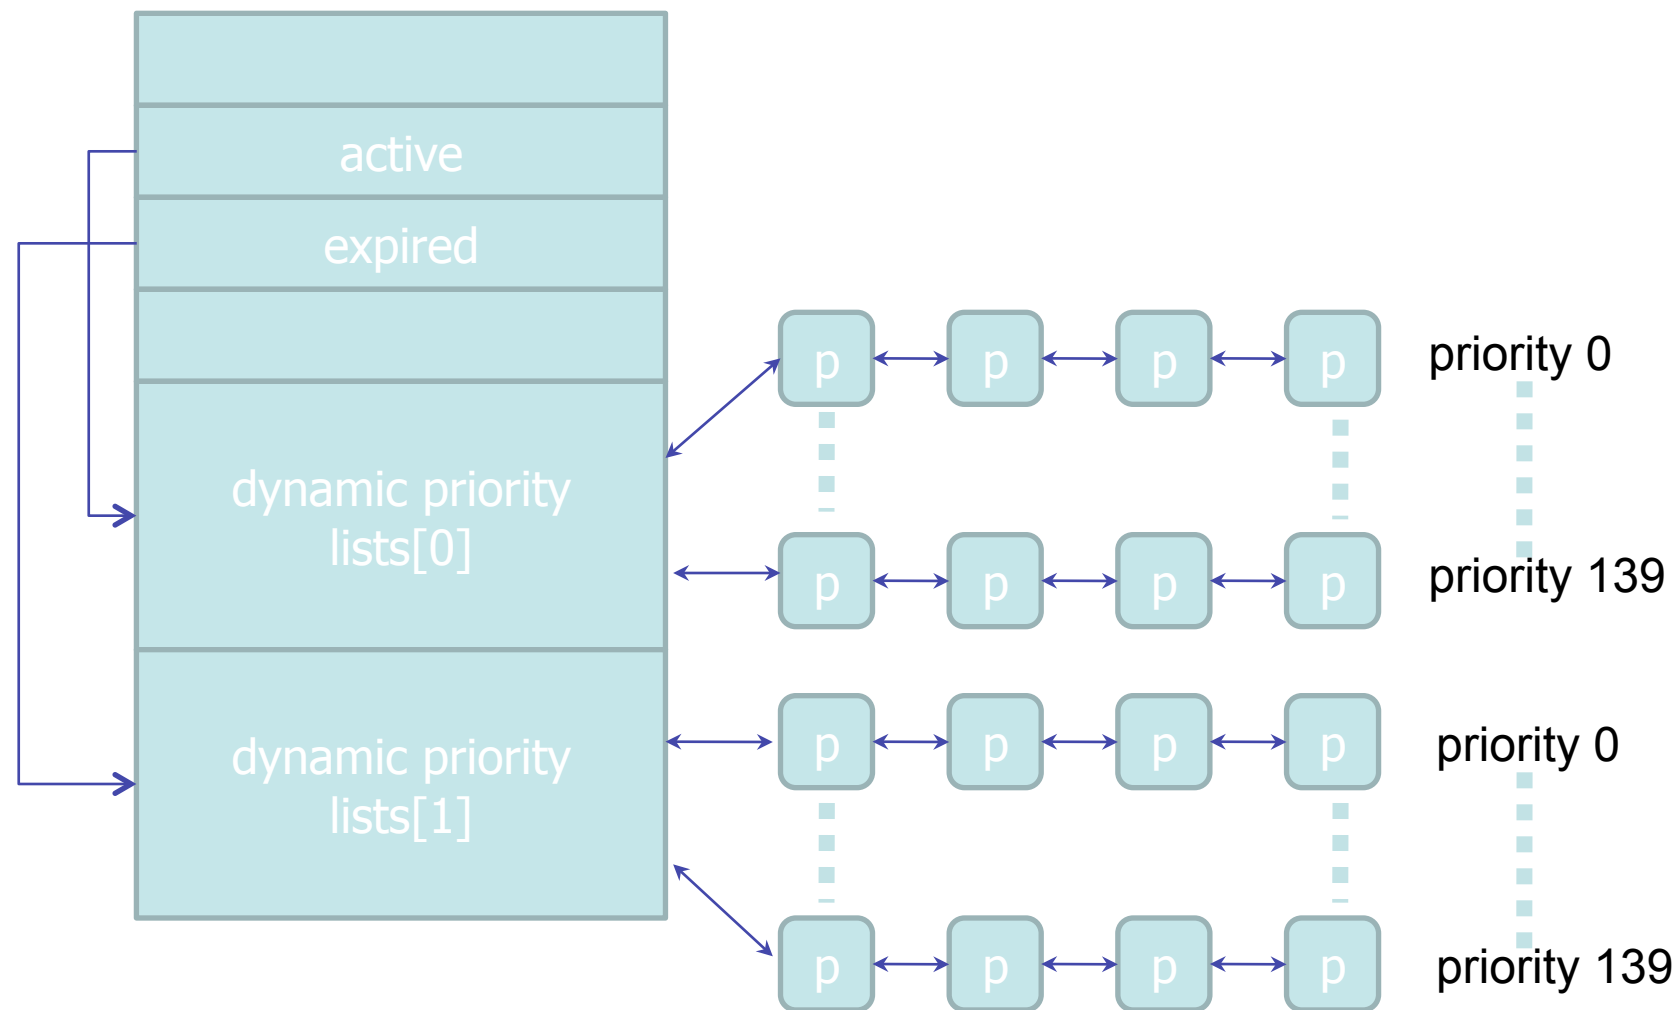


Process scheduling in Linux

- Real-time process scheduling
 - (Soft) real-time processes are assigned priorities in the range 1 (highest priority) ... 99 (lowest priority)
 - There are two classes of real-time process: FIFO and “round-robin”
 - The priority is not dynamic
- FIFO processes
 - continue executing until they block or until they are pre-empted by a higher priority process
 - time slice is irrelevant
- “Round-Robin” processes
 - continue executing until they block, are pre-empted or their time slice expires
 - when the time-slice of a Round-Robin process expires, it is placed at the tail of the priority list



Process scheduling in Linux



Process scheduling in Linux

- Periodically, the remaining time slice of the currently executing process is decremented
- If the time slice has been exhausted ...
 - the process is removed from the active queue
 - its time slice is restored
- ... and it is placed ...
 - on the active list if it is interactive and there are no starving expired processes
 - on the expired list otherwise
- Schedule the next process to execute from the head of the highest priority active queue
- If all processes are expired, the expired and active lists are swapped (cheaply, by changing pointers)

