

Concurrent Systems Operating Systems

3D4  CS2016

Andrew Butterfield
ORI.G39, Andrew.Butterfield@scss.tcd.ie



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

with thanks to Mike Brady

The xv6 operating system

- Developed by MIT for teaching
 - Keeps it simple, rather than efficient!
 - <https://pdos.csail.mit.edu/6.828/2012/xv6.html> (or Google 'xv6')
- Open source
 - github.com/mit-pdos/xv6-public.git
- Code Listing
 - <https://pdos.csail.mit.edu/6.828/2018/xv6/xv6-rev11.pdf>
- Accompanying 'Book'
 - <https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev11.pdf>



The xv6 'book'

- V. Good introduction to OS design - Chapter 0 is well worth a read!
- Each chapter handles a topic: problem, solution, code commentary

The function `switch` performs the saves and restores for a thread switch. `switch` doesn't directly know about threads; it just saves and restores register sets, called *contexts*. When it is time for a process to give up the CPU, the process's kernel thread calls `switch` to save its own context and return to the scheduler context. Each context is represented by a `struct context*`, a pointer to a structure stored on the kernel stack involved. `Switch` takes two arguments: `struct context **old` and `struct context *new`. It pushes the current registers onto the stack and saves the stack pointer in `*old`. Then `switch` copies `new` to `%esp`, pops previously saved registers, and returns.

Let's follow a user process through `switch` into the scheduler. We saw in Chapter 3 that one possibility at the end of each interrupt is that `trap` calls `yield`. `Yield` in turn calls `sched`, which calls `switch` to save the current context in `proc->context` and switch to the scheduler context previously saved in `cpu->scheduler` (2822).



The xv6 code-listing

```
2800 // Enter scheduler. Must hold only ptable.lock
2801 // and have changed proc->state. Saves and restores
2802 // intena because intena is a property of this
2803 // kernel thread, not this CPU. It should
2804 // be proc->intena and proc->ncli, but that would
2805 // break in the few places where a lock is held but
2806 // there's no process.
2807 void
2808 sched(void)
2809 {
2810     int intena;
2811     struct proc *p = myproc();
2812
2813     if(!holding(&ptable.lock))
2814         panic("sched ptable.lock");
2815     if(mycpu()->ncli != 1)
2816         panic("sched locks");
2817     if(p->state == RUNNING)
2818         panic("sched running");
2819     if(readeflags() & FL_IF)
2820         panic("sched interruptible");
2821     intena = mycpu()->intena;
2822     swtch(&p->context, mycpu()->scheduler);
2823     mycpu()->intena = intena;
2824 }
```

- Smart way to do listing:
 - 2-columns of 50 lines, so line-numbers have page number
- Each chapter handles a topic:
 - problem,
 - solution,
 - code commentary



The xv6 scheduler (I)

- Multicore!
 - we show an overview here (inner loop body not shown ‘...’)
- Key variables:
 - Pointer `c` will point to selected process
 - Pointer `p` is used to walk the process table (`ptable`) from start to finish
 - Note the use of locking for the process table!
- Book p63, Listings 2750-2791

```
// Per-CPU process scheduler.
// Each CPU calls scheduler() after setting itself up.
// Scheduler never returns. It loops, doing:
//  - choose a process to run
//  - swtch to start running that process
//  - eventually that process transfers control
//    via swtch back to the scheduler.
void
scheduler(void)
{ struct proc *p;
  struct cpu *c = mycpu(); c->proc = 0;
  for(;;){
    sti(); // Enable interrupts on this processor.
    // Loop over process table looking for process to run.
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
      ...
    }
    release(&ptable.lock);
  }
}
```



The xv6 scheduler (2)

- We skip non-runnable processes
 - We choose the next runnable process
 - Round-Robin scheduler
 - Simple
- Prep. for running it:
 - Mark as running, and set processor mode to 'user' (switchvm)
- Running it: call `swtch`
- When process returns to scheduler:
 - set processor mode to 'kernel'
 - set selected process pointer to zero.

```
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->state != RUNNABLE)
        continue;

    // Switch to chosen process. It is the process's job
    // to release ptable.lock and then reacquire it
    // before jumping back to us.
    c->proc = p;
    switchvm(p);
    p->state = RUNNING;

    swtch(&(c->scheduler), p->context);
    switchkvm();

    // Process is done running for now.
    // It should have changed its p->state before coming back.
    c->proc = 0;
}
```



The xv6 process table (I)

- Table has a 'spinlock' used to manage access
- It has an array of struct proc
 - contains 'metadata' for each process.
- Lots of headers
 - At a guess, perhaps proc.h is relevant?
- Listings 2400-2412

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "x86.h"
#include "proc.h"
#include "spinlock.h"

struct {
    struct spinlock lock;
    struct proc proc[NPROC];
} ptable;
```



The xv6 process table (II)

- Complicated!
- Most not relevant to scheduling, though.
- Relevant :
sz pid foil cwd
(?)
- Listings 2334-2351

```
// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;               // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;            // If non-zero, sleeping on chan
    int killed;            // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;     // Current directory
    char name[16];         // Process name (debugging)
};
```

```
enum procstate { UNUSED, EMBRYO, SLEEPING
                , RUNNABLE, RUNNING, ZOMBIE };
```



The xv6 process context

- A C struct containing the relevant 6 registers, exactly as they appear on the stack
- Listings 2316-2332

```
// Saved registers for kernel context switches.  
// Don't need to save all the segment registers (%cs, etc),  
// because they are constant across kernel contexts.  
// Don't need to save %eax, %ecx, %edx, because the  
// x86 convention is that the caller has saved them.  
// Contexts are stored at the bottom of the stack they  
// describe; the stack pointer is the address of the context.  
// The layout of the context matches the layout of the stack in swtch.S  
// at the "Switch stacks" comment. Switch doesn't save eip explicitly,  
// but it is on the stack and allocproc() manipulates it.  
struct context {  
    uint edi;  
    uint esi;  
    uint ebx;  
    uint ebp;  
    uint eip;  
};
```



xv6 context switching

- Assembly Language!
 - x86
- Switches out a process or the kernel
- Switches in the kernel or a process
- A C prototype is provided
- Listings 3050-3078

```
# Context switch
#
# void switch(struct context **old, struct context *new);
#
# Save the current registers on the stack, creating
# a struct context, and save its address in *old.
# Switch stacks to new and pop previously-saved registers.
```

```
.globl switch
switch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx
```

```
# Save old callee-saved registers
pushl %ebp
pushl %ebx
pushl %esi
pushl %edi
```

```
# Switch stacks
movl %esp, (%eax)
movl %edx, %esp
```

```
# Load new callee-saved registers
popl %edi
popl %esi
popl %ebx
popl %ebp
ret
```



xv6 and CS2016/3D4

- Full use of xv6 requires
 - installing an emulator called Qemu, with MIT 'tweaks'
 - being proficient with the gdb debugger
 - installing cross-compilers if not using Unix
 - building familiarity over time
- Our use:
 - we shall abstract out the scheduler
 - provide a test harness program to exercise it
 - so you can improve the simple algorithm

