

# Concurrent Systems Operating Systems

3D4 ← → CS2016

*Andrew Butterfield*  
*ORI.G39, [Andrew.Butterfield@scss.tcd.ie](mailto:Andrew.Butterfield@scss.tcd.ie)*



**Trinity College Dublin**  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

*with thanks to Mike Brady*

# recent OS research @SCSS

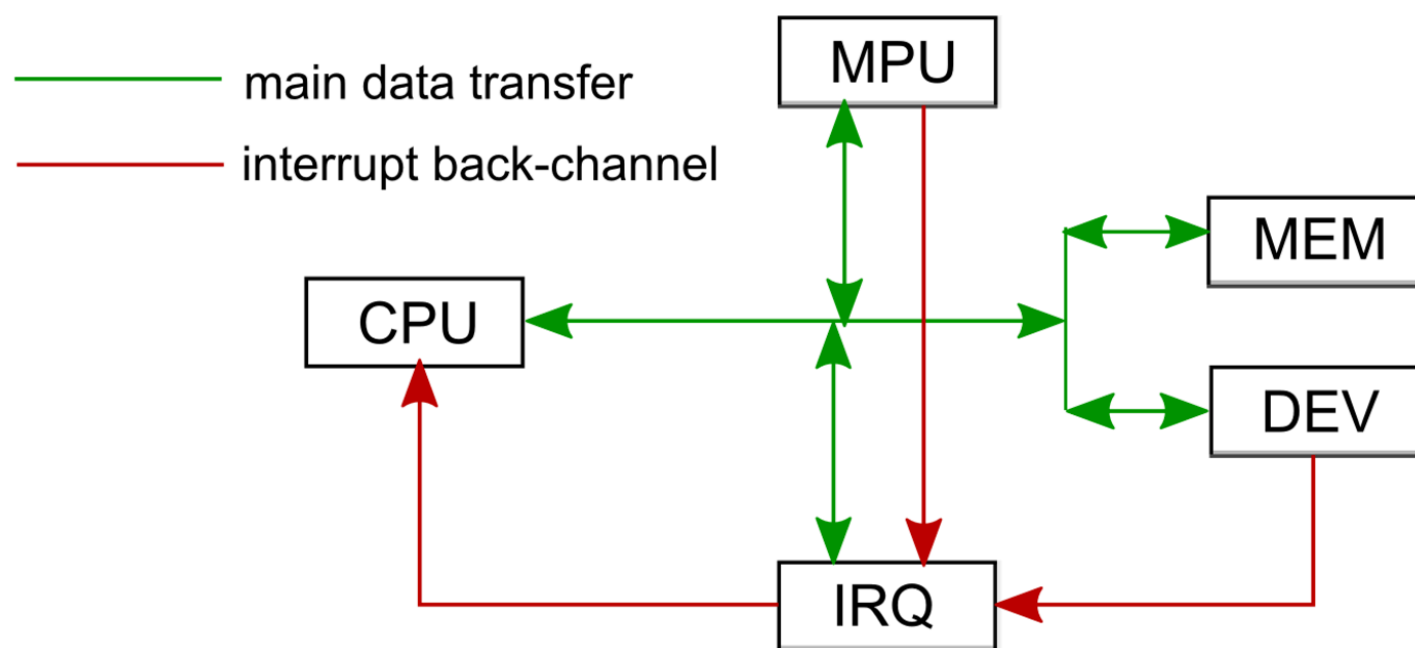
- Applying formal methods to OS verification
- Why?
  - Testing is king, but requires test-setup repeatability.
  - OS behaviour is effectively non-deterministic:
    - mainly because of interrupt unpredictability
    - hard to repeat tests
    - Multicore makes things worse



# From FMEIMAKQP(?) final presentation

## Unavoidable Concurrency – even with Single-Core!

- In a single-core system, the CPU is time-shared between the hypervisor and partition code with no parallelism.
- Flow control change is managed via traps (exceptions, interrupts, ...).
- However, we have an essentially concurrent system
- CPU executes instructions in a sequential manner on behalf of either the kernel or partition
- Memory (MEM) responds to CPU memory requests
- The MMU/MPU observes the bus traffic, raises memory fault interrupts when appropriate
- IO Devices (DEV) signal via interrupt when done
- Interrupt request hardware (IRQ) takes in interrupt requests from MMU/MPU/DEV/CPU and forwards the highest in priority to the CPU



# Formal Methods?

- Use of mathematical logics to:
  - Model application domains
  - Give precise semantics to requirements, specification and programming languages
  - Allow proofs of program correctness
  - Provide a rigorous foundation for developing analysis tools
    - e.g., model-checkers, like Promela/SPIN.



# What kinds of OS?

- Real-time OSes used in spacecraft
  - research/consultancy for European Space Agency
- Time/Space Partitioning (TSP) Kernels
  - a form of “hypervisor”
  - two projects: MTOBSE, FMEIMAKQP
- Real-Time Executive for Multiprocessor Systems (RTEMS)
  - project starting: RTEMS-SMP

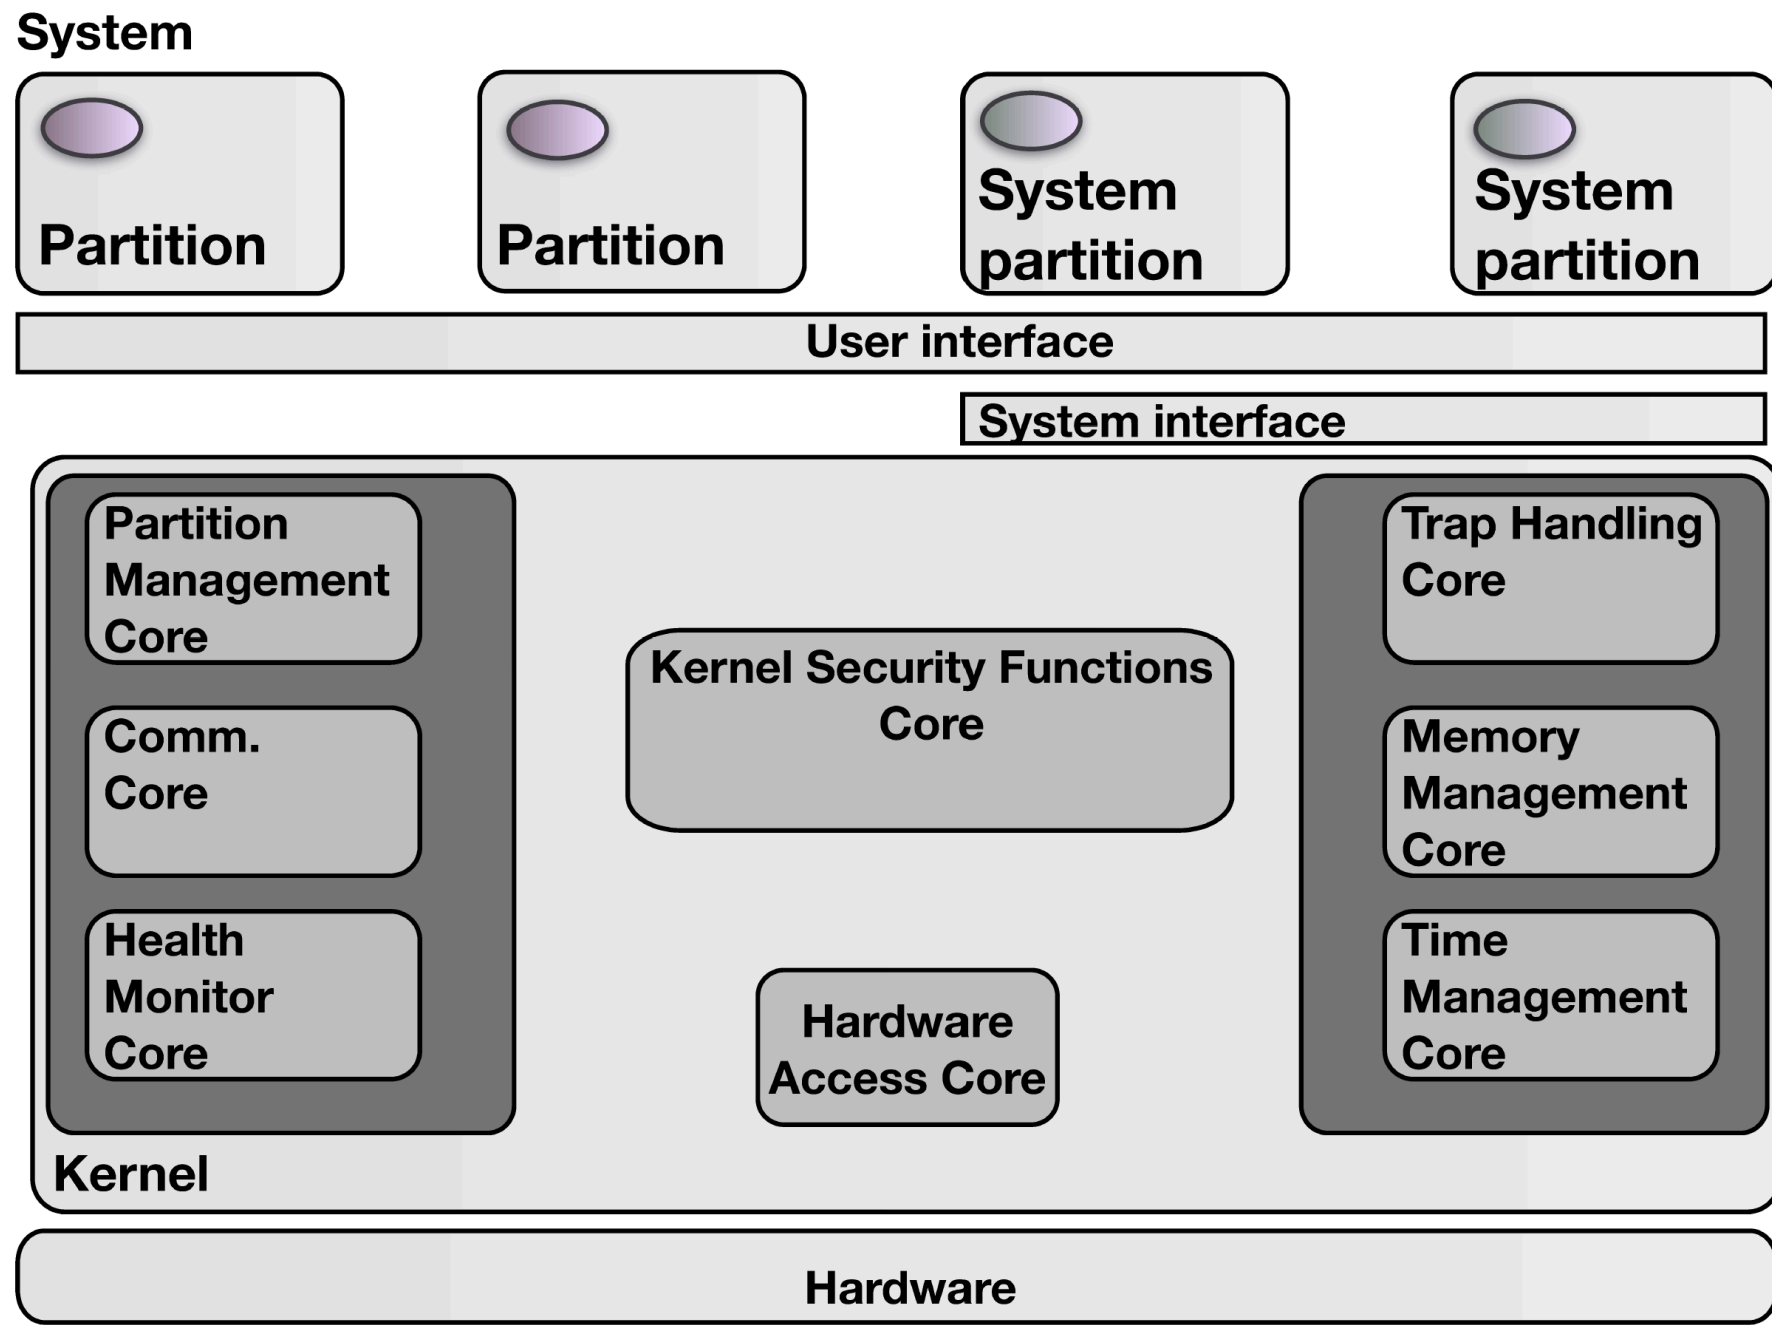


# Lero - the Irish Software Research Centre

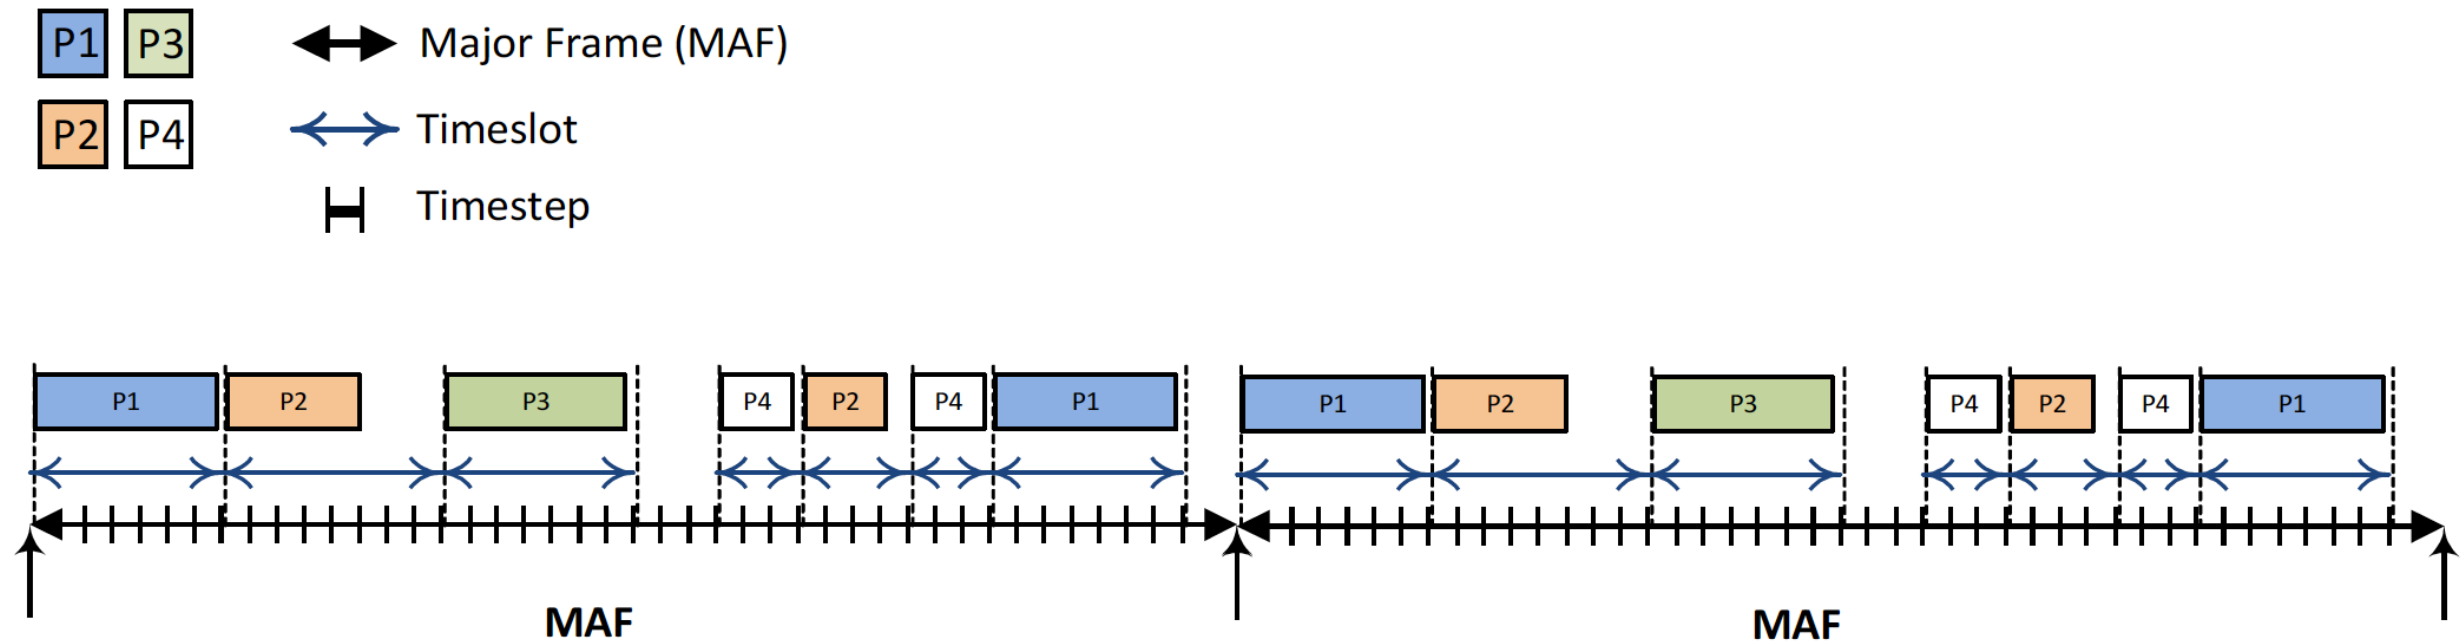
- SFI-funded Research Centre
- Involves all 7 universities in Ireland, as well as some of the IT sector
  - Headquarters in Limerick
- Strong industrial focus
- Main point of contact with ESA



# Time/Space Partitioning Kernel



# Fixed TSP Scheduling



“Schedulability” analysis is crucial in order to work out the fixed timings above!

IMA Separation Kernel Qualification Project

ESA Contract N°: 4000111495/14/NL/GLC/al

**Do2: Partitioning Kernel Requirements Baseline**



Trinity College Dublin  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin



# Excerpts from an TSP formal model

```
type_synonym tfwWLength= nat
type_synonym partitionID = nat
type_synonym tfwOffset = nat
type_synonym timeFrameWindow = "(partitionID * tfwWLength * tfwOffset)"
type_synonym majorTimeFrame = "timeFrameWindow list"
type_synonym TFWid = "nat"
datatype schedule = Sch TFWid majorTimeFrame
...
definition
  scheduler :: "unit s_monad"
where
  "scheduler  $\equiv$  do
    parMan  $\leftarrow$  gets partitionManager;
    schedule  $\leftarrow$  return (partitionSchedule parMan);
    nextTFWNonIdle  $\leftarrow$  nextTFW(schedule);
    schedule'  $\leftarrow$  setTFWSchedule schedule nextTFWNonIdle;
    modify ( $\lambda s . s$ (partitionManager := (partitionManager s)(partitionSchedule := schedule')));
    s  $\leftarrow$  get;
    archSwitchToPartition
      getCurrentTFW(partitionSchedule(partitionManager s))
  od"
...
lemma sanitization:
 $\exists s, s' \in \text{state} . s' = \text{scheduler } s \Rightarrow \forall r \in (\text{cpu\_regs } (\text{cpu\_state } s')) . r=0$ 
by (simp_all add:archSwitchToPartition add:scheduler)
```



# Using CSP to model hardware

## CSP Model of MMU

```
MMU(blocked)
=  tick -> ( TRANSFER(blocked)
              [] BLOCK_MEM_REGION(blocked)
              [] UNBLOCK_MEM_REGION(blocked)
              [] MMU(blocked) )

TRANSFER(blocked)
=  bus?dir?addr -> ( if member(addr, blocked)
                     then ( badaccess -> raise!memfault -> MMU(blocked) )
                     else ( mmuOK -> MMU(blocked) )
                   )

BLOCK_MEM_REGION(blocked)
=  mmuBlock?a -> MMU(union(blocked,{a}))

UNBLOCK_MEM_REGION(blocked)
=  mmuUnBlock?a -> MMU(diff(blocked,{a}))
```

- Synchronise with global clock (**tick**)
- Accepts bus read or write (**bus?dir?addr**)
- Permits operation if address not blocked (**blocked**, **mmuOK**)
- Objects if address is blocked (**blocked**, **badaccess**, **raise!memfault**)

CSP: Communicating Sequential Processes, developed by C.A.R. Hoare

Has powerful model-checker called FDR (Failures-Divergences Refinement)



# RTEMS, “improved”

- RTEMS: open-source real-time operating system ([rtems.org](http://rtems.org))
  - Widely used
- RTEMS Improvement
  - A version of single-core RTEMS qualified for spaceflight
    - done for ESA by Thales Edisoft (Portugal)
  - Code in RTEMS repository, qualification data owned by Edisoft.



# RTEMS-SMP

- In 2013, two good realtime scheduling algorithms for multicore emerged
  - Multiprocessor resource sharing Protocol (MrsP)
    - A. Burns and A.J. Wellings, U. of York
  - O(m) independence preserving protocol (OMIP)
    - Bjorn B. Brandenburg, Max Planck Institute for Software Systems
- Within a few years, an RTEMS version was written by an employee of a German company, Embedded Brains
  - called RTEMS-SMP, part of the RTEMS repo, to be in the next major release.



# RTEMS-SMP pre-qualification

- ESA want a flight qualified version of RTEMS-SMP
- Edisoft and Embedded Brains, along with a German instrument maker called Jena Optronik formed a consortium to do this.
- ESA invited key formal methods researchers from [Lero](#)
  - Andrew Butterfield Lero@TCD
  - Mike Hinchey Lero@UL, also NASA consultant.
- 2-year project, started mid-Feb 2019



# Formal Aspects of RTEMS-SMP

- Main focus: the use of MrsP and OMIP in RTEMS
  - static analysis tools (coverity, infer,...)
  - model checking (SPIN,TLC, FDR)
  - Correctness Proofs (Frama-C,TLA+, ...)

